

# Model Evaluation

October 1, 2024

## 0.1 Data Checking

```
[1]: import pandas as pd
import numpy as np
```

```
[2]: data = pd.read_csv('Flyzy Flight Cancellation - Sheet1.csv')
data.iloc[1:4, 2:3]= np.NaN
data.iloc[1:4, 3:4]= "NA"
data.iloc[1:4, 4:5]= ""
data["None_col"]= None
data.head()
```

```
[2]:
```

	Flight ID	Airline	Flight_Distance	Origin_Airport	Destination_Airport	\
0	7319483	Airline D	475.0	Airport 3	Airport 2	
1	4791965	Airline E	NaN	NA		
2	2991718	Airline C	NaN	NA		
3	4220106	Airline E	NaN	NA		
4	2263008	Airline E	566.0	Airport 2	Airport 2	

	Scheduled_Departure_Time	Day_of_Week	Month	Airplane_Type	Weather_Score	\
0	4	6	1	Type C	0.225122	
1	12	1	6	Type B	0.060346	
2	17	3	9	Type C	0.093920	
3	1	1	8	Type B	0.656750	
4	19	7	12	Type E	0.505211	

	Previous_Flight_Delay_Minutes	Airline_Rating	Passenger_Load	\
0	5.0	2.151974	0.477202	
1	68.0	1.600779	0.159718	
2	18.0	4.406848	0.256803	
3	13.0	0.998757	0.504077	
4	4.0	3.806206	0.019638	

	Flight_Cancelled	None_col
0	0	None
1	1	None
2	0	None
3	1	None

4                      0        None

```
[3]: null= pd.isnull(data)
      null.head()
```

```
[3]:
```

	Flight ID	Airline	Flight_Distance	Origin_Airport	Destination_Airport	\
0	False	False	False	False	False	
1	False	False	True	False	False	
2	False	False	True	False	False	
3	False	False	True	False	False	
4	False	False	False	False	False	

	Scheduled_Departure_Time	Day_of_Week	Month	Airplane_Type	Weather_Score	\
0	False	False	False	False	False	
1	False	False	False	False	False	
2	False	False	False	False	False	
3	False	False	False	False	False	
4	False	False	False	False	False	

	Previous_Flight_Delay_Minutes	Airline_Rating	Passenger_Load	\
0	False	False	False	
1	False	False	False	
2	False	False	False	
3	False	False	False	
4	False	False	False	

	Flight_Cancelled	None_col
0	False	True
1	False	True
2	False	True
3	False	True
4	False	True

```
[4]: pd.isnull(data).sum().sum()
```

```
[4]: 3003
```

```
[5]: missing_values = data.isnull().sum()
      print("Missing values per column (before handling):")
      print(missing_values)
```

```
Missing values per column (before handling):
Flight ID                0
Airline                  0
Flight_Distance          3
Origin_Airport           0
Destination_Airport      0
```

```

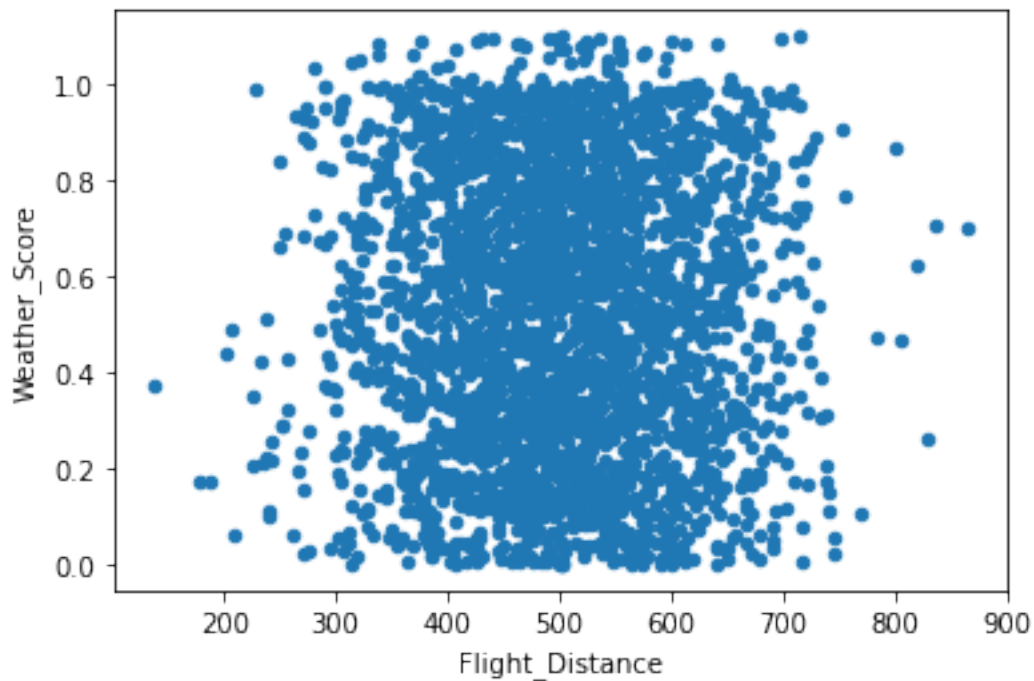
Scheduled_Departure_Time      0
Day_of_Week                   0
Month                         0
Airplane_Type                 0
Weather_Score                 0
Previous_Flight_Delay_Minutes 0
Airline_Rating                0
Passenger_Load                0
Flight_Cancelled              0
None_col                      3000
dtype: int64

```

```

[6]: import pandas as pd
import matplotlib.pyplot as plt
data.plot(kind='scatter' , x= 'Flight_Distance', y= 'Weather_Score')
plt.show()

```




```

[7]: import matplotlib.pyplot as plt
import pandas as pd
columns_to_check = ['Flight_Distance', 'Weather_Score']
plt.figure(figsize=(10, 6))
data[columns_to_check].plot(kind= 'box')
plt.title('Boxplots for Outlier Detection')
plt.xlabel('Features')
plt.ylabel('Values')

```

```
plt.grid(True)
plt.show()
```

<Figure size 720x432 with 0 Axes>



```
[8]: data_types = data.dtypes
print("Data types of each column:")
print(data_types)
```

Data types of each column:

Flight ID	int64
Airline	object
Flight_Distance	float64
Origin_Airport	object
Destination_Airport	object
Scheduled_Departure_Time	int64
Day_of_Week	int64
Month	int64
Airplane_Type	object
Weather_Score	float64
Previous_Flight_Delay_Minutes	float64
Airline_Rating	float64
Passenger_Load	float64

```
Flight_Cancelled          int64
None_col                  object
dtype: object
```

```
[9]: Data =pd.read_csv('Flyzy Flight Cancellation - Sheet1.csv')
```

```
[10]: print(Data.head())
```

	Flight ID	Airline	Flight_Distance	Origin_Airport	Destination_Airport	\
0	7319483	Airline D	475	Airport 3	Airport 2	
1	4791965	Airline E	538	Airport 5	Airport 4	
2	2991718	Airline C	565	Airport 1	Airport 2	
3	4220106	Airline E	658	Airport 5	Airport 3	
4	2263008	Airline E	566	Airport 2	Airport 2	

	Scheduled_Departure_Time	Day_of_Week	Month	Airplane_Type	Weather_Score	\
0		4	6	1	Type C	0.225122
1		12	1	6	Type B	0.060346
2		17	3	9	Type C	0.093920
3		1	1	8	Type B	0.656750
4		19	7	12	Type E	0.505211

	Previous_Flight_Delay_Minutes	Airline_Rating	Passenger_Load	\
0	5.0	2.151974	0.477202	
1	68.0	1.600779	0.159718	
2	18.0	4.406848	0.256803	
3	13.0	0.998757	0.504077	
4	4.0	3.806206	0.019638	

	Flight_Cancelled
0	0
1	1
2	0
3	1
4	0

```
[11]: print(Data.head())
```

	Flight ID	Airline	Flight_Distance	Origin_Airport	Destination_Airport	\
0	7319483	Airline D	475	Airport 3	Airport 2	
1	4791965	Airline E	538	Airport 5	Airport 4	
2	2991718	Airline C	565	Airport 1	Airport 2	
3	4220106	Airline E	658	Airport 5	Airport 3	
4	2263008	Airline E	566	Airport 2	Airport 2	

	Scheduled_Departure_Time	Day_of_Week	Month	Airplane_Type	Weather_Score	\
0		4	6	1	Type C	0.225122
1		12	1	6	Type B	0.060346

2	17	3	9	Type C	0.093920
3	1	1	8	Type B	0.656750
4	19	7	12	Type E	0.505211

	Previous_Flight_Delay_Minutes	Airline_Rating	Passenger_Load \
0	5.0	2.151974	0.477202
1	68.0	1.600779	0.159718
2	18.0	4.406848	0.256803
3	13.0	0.998757	0.504077
4	4.0	3.806206	0.019638

	Flight_Cancelled
0	0
1	1
2	0
3	1
4	0

## EDA ##

```
[12]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[13]: Data =pd.read_csv('Flyzy Flight Cancellation - Sheet1.csv')
```

```
[14]: print(Data.head())
```

	Flight ID	Airline	Flight_Distance	Origin_Airport	Destination_Airport \
0	7319483	Airline D	475	Airport 3	Airport 2
1	4791965	Airline E	538	Airport 5	Airport 4
2	2991718	Airline C	565	Airport 1	Airport 2
3	4220106	Airline E	658	Airport 5	Airport 3
4	2263008	Airline E	566	Airport 2	Airport 2

	Scheduled_Departure_Time	Day_of_Week	Month	Airplane_Type	Weather_Score \
0	4	6	1	Type C	0.225122
1	12	1	6	Type B	0.060346
2	17	3	9	Type C	0.093920
3	1	1	8	Type B	0.656750
4	19	7	12	Type E	0.505211

	Previous_Flight_Delay_Minutes	Airline_Rating	Passenger_Load \
0	5.0	2.151974	0.477202
1	68.0	1.600779	0.159718
2	18.0	4.406848	0.256803
3	13.0	0.998757	0.504077

4	4.0	3.806206	0.019638
---	-----	----------	----------

Flight_Cancelled	
0	0
1	1
2	0
3	1
4	0


```
[15]: print(data.isnull().sum())
```

Flight ID	0
Airline	0
Flight_Distance	3
Origin_Airport	0
Destination_Airport	0
Scheduled_Departure_Time	0
Day_of_Week	0
Month	0
Airplane_Type	0
Weather_Score	0
Previous_Flight_Delay_Minutes	0
Airline_Rating	0
Passenger_Load	0
Flight_Cancelled	0
None_col	3000
dtype:	int64

```
[16]: print(Data.dtypes)
```

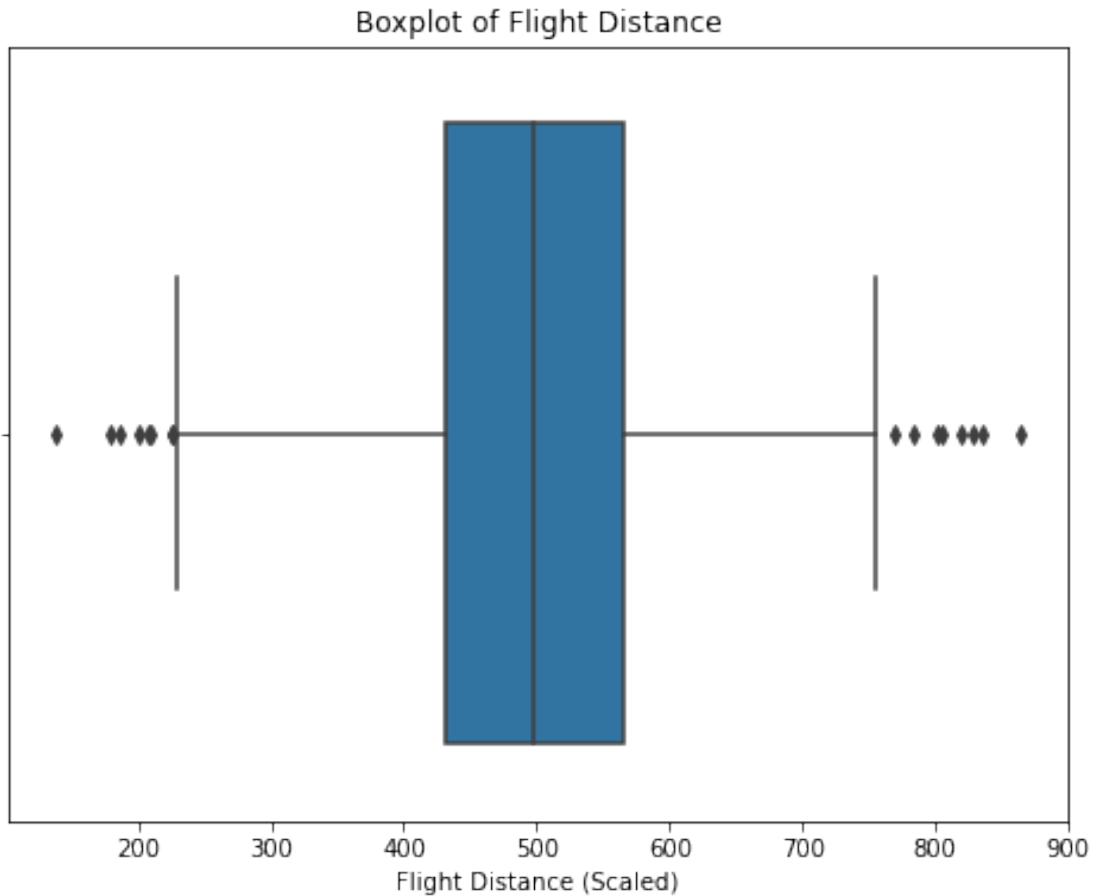
Flight ID	int64
Airline	object
Flight_Distance	int64
Origin_Airport	object
Destination_Airport	object
Scheduled_Departure_Time	int64
Day_of_Week	int64
Month	int64
Airplane_Type	object
Weather_Score	float64
Previous_Flight_Delay_Minutes	float64
Airline_Rating	float64
Passenger_Load	float64
Flight_Cancelled	int64
dtype:	object

```
[17]: plt.figure(figsize=(12,8))
sns.boxplot(data=Data)
plt.xticks(rotation=45)
plt.show()
```

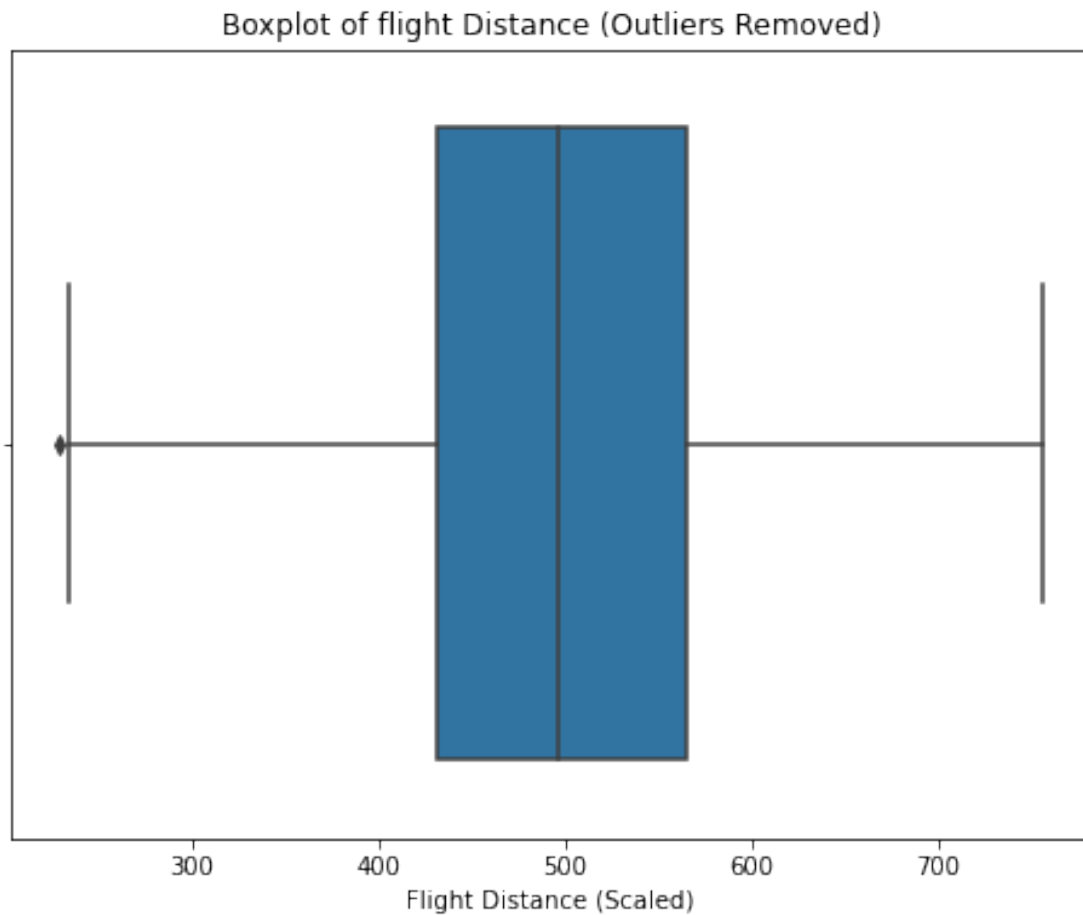


```
[18]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Flight_Distance'])
plt.title('Boxplot of Flight Distance')
plt.xlabel('Flight Distance (Scaled)')
plt.show()
```






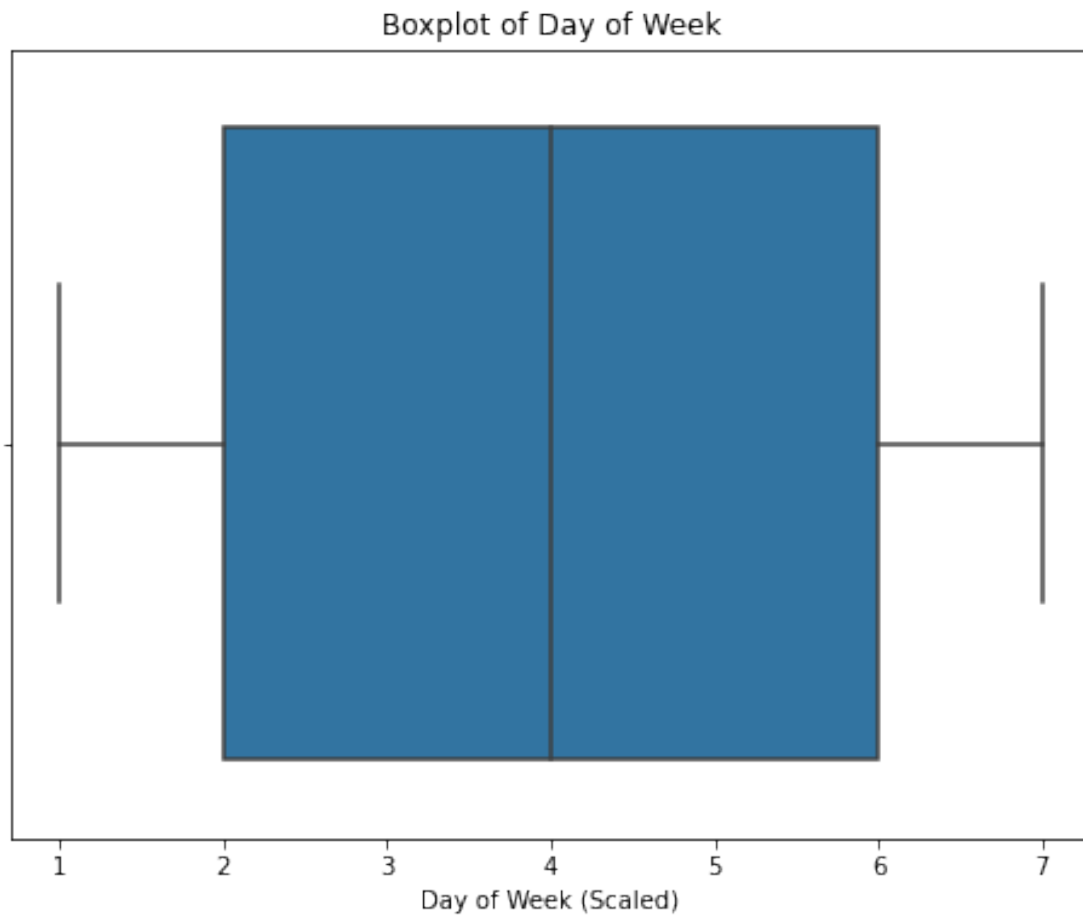
```
[19]: Q1 = Data['Flight_Distance'].quantile(0.25)
      Q3 = Data['Flight_Distance'].quantile(0.75)
      IQR = Q3 - Q1
      lower_bound = Q1 - 1.5 * IQR
      upper_bound = Q3 + 1.5 * IQR
      filtered_Data = Data[(Data['Flight_Distance'] >= lower_bound) &
                           (Data['Flight_Distance'] <= upper_bound)]
      plt.figure(figsize=(8, 6))
      sns.boxplot(x=filtered_Data['Flight_Distance'])
      plt.title('Boxplot of flight Distance (Outliers Removed)')
      plt.xlabel('Flight Distance (Scaled)')
      plt.show()
```



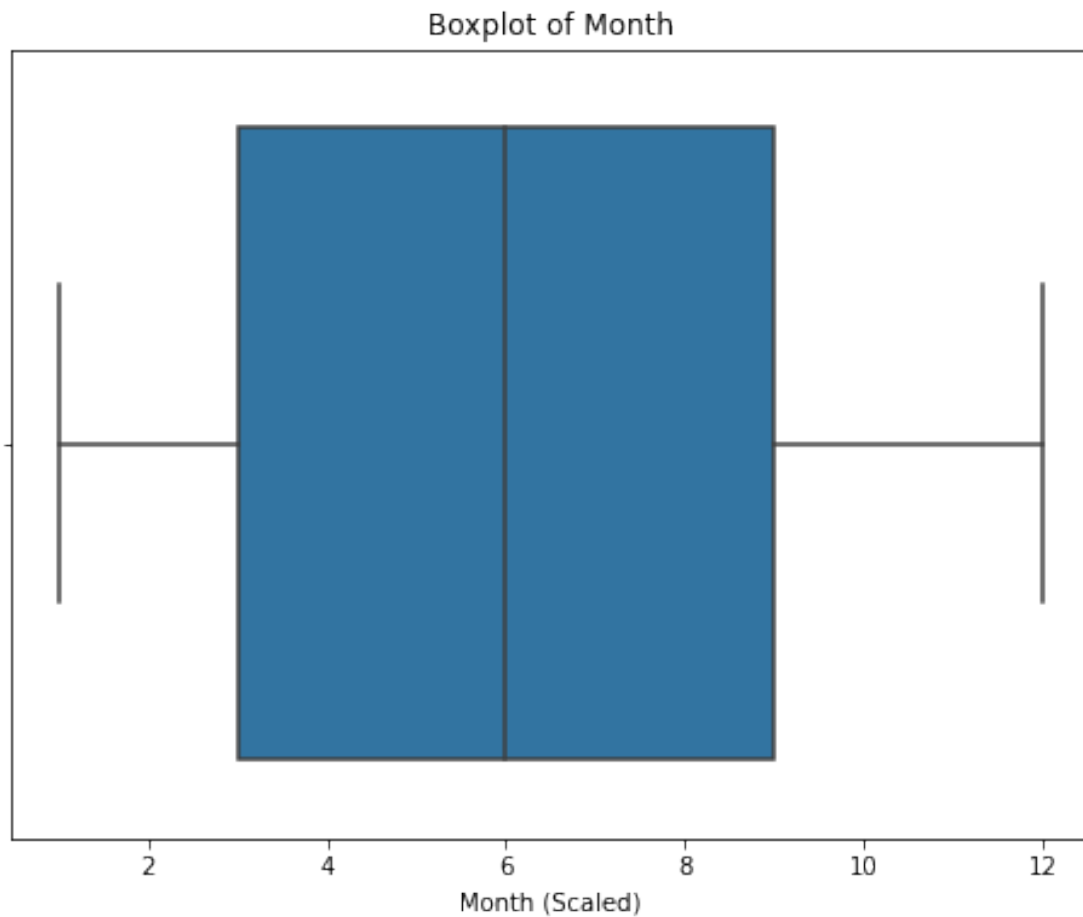
```
[20]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Scheduled_Departure_Time'])
plt.title('Boxplot of Scheduled Departure Time')
plt.xlabel('Scheduled Departure Time (Scaled)')
plt.show()
```



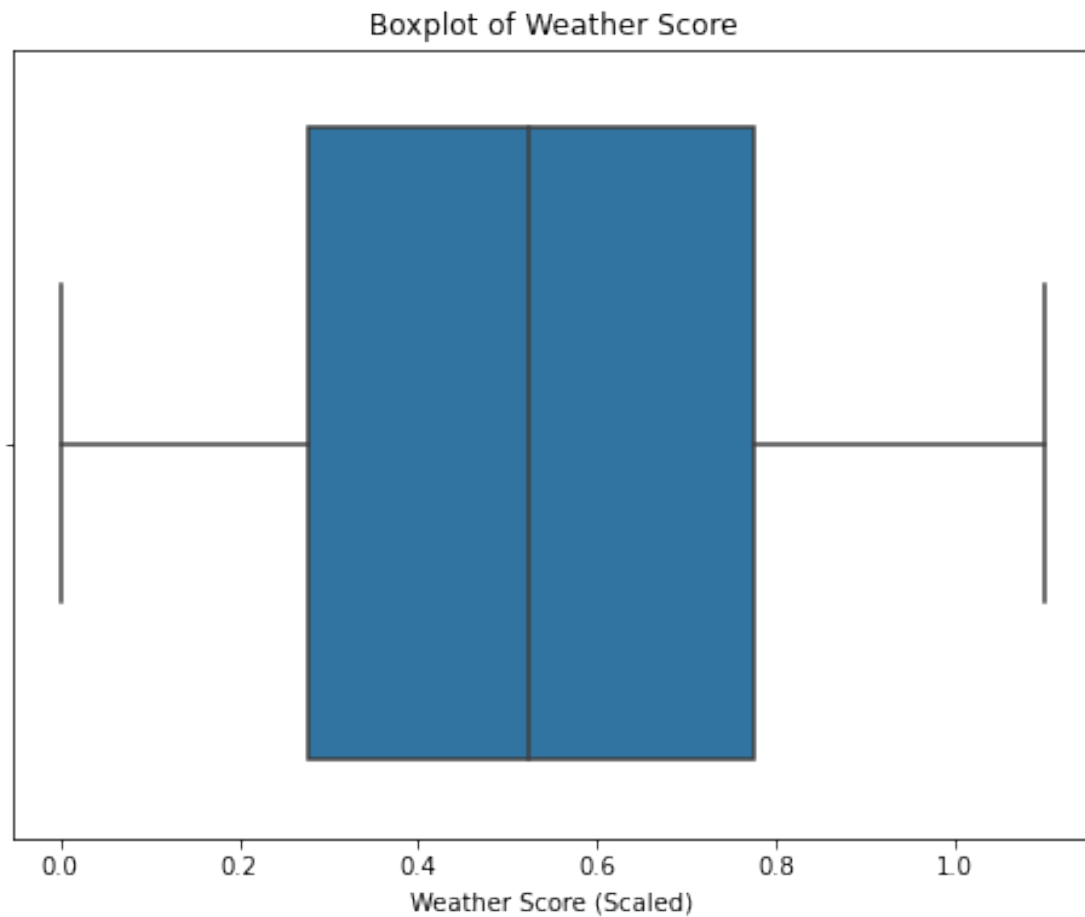
```
[21]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Day_of_Week'])
plt.title('Boxplot of Day of Week')
plt.xlabel('Day of Week (Scaled)')
plt.show()
```



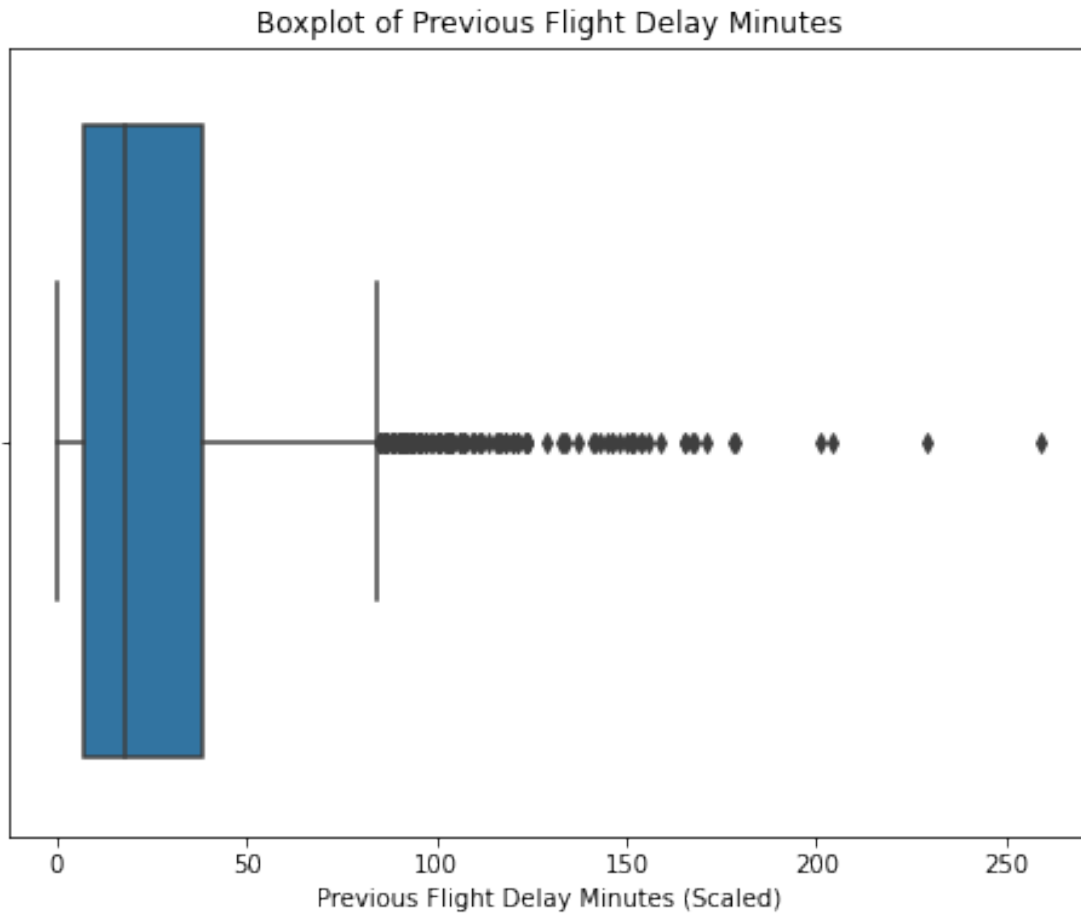
```
[22]: plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Month'])
plt.title('Boxplot of Month')
plt.xlabel('Month (Scaled)')
plt.show()
```



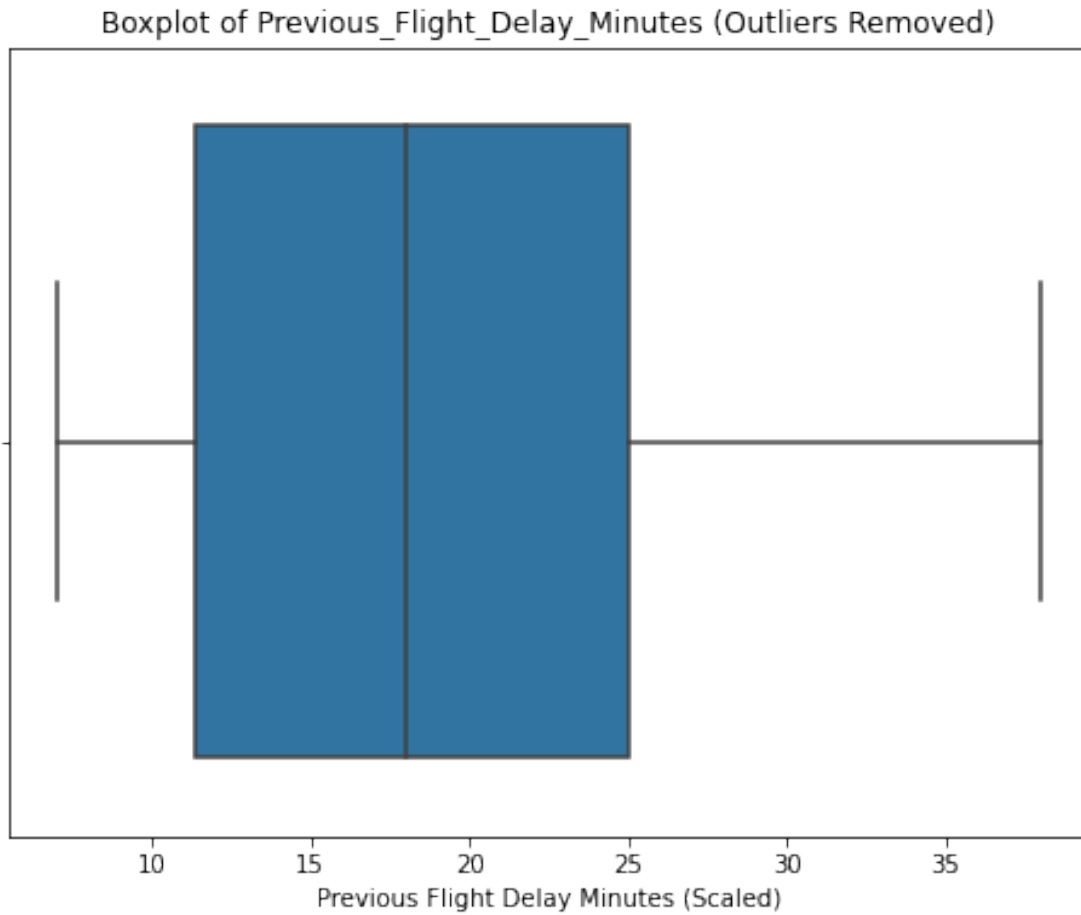
```
[23]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Weather_Score'])
plt.title('Boxplot of Weather Score')
plt.xlabel('Weather Score (Scaled)')
plt.show()
```



```
[24]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Previous_Flight_Delay_Minutes'])
plt.title('Boxplot of Previous Flight Delay Minutes')
plt.xlabel('Previous Flight Delay Minutes (Scaled)')
plt.show()
```

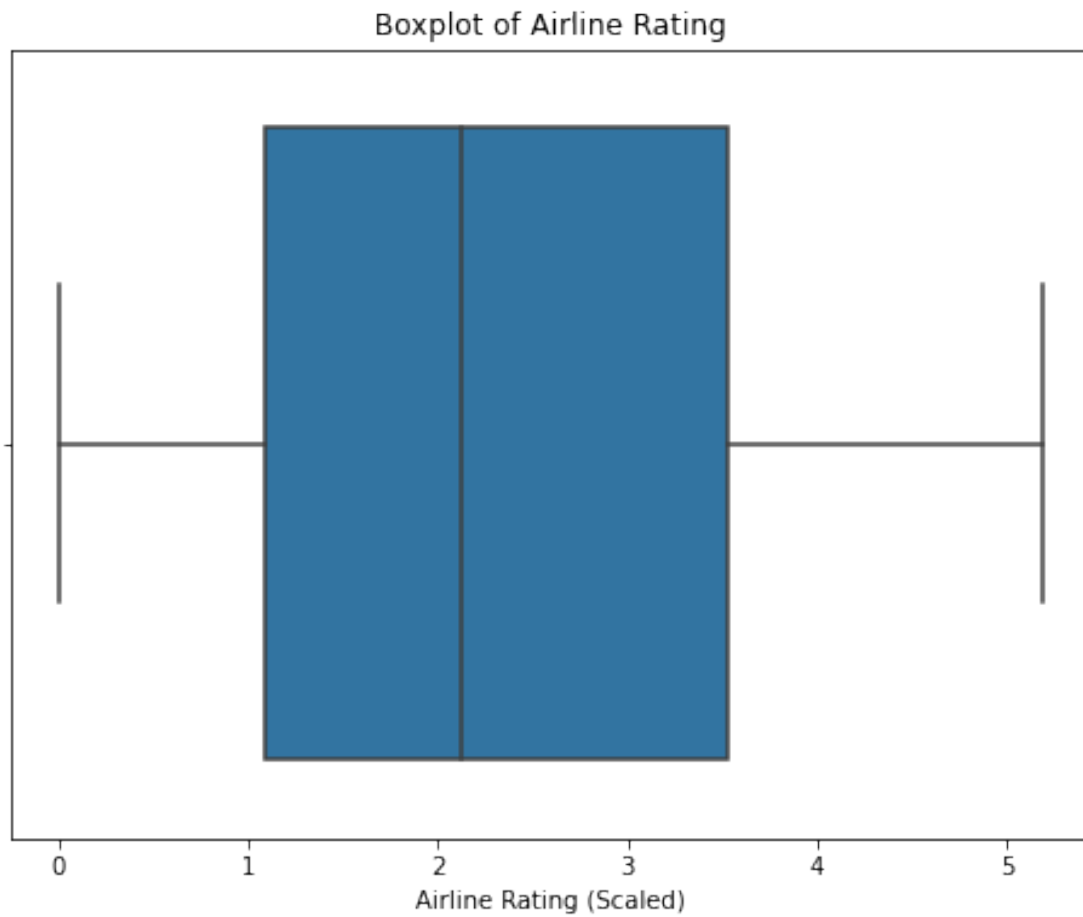


```
[25]: Q1 = Data['Previous_Flight_Delay_Minutes'].quantile(0.25)
Q3 = Data['Previous_Flight_Delay_Minutes'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 0.0 * IQR
upper_bound = Q3 + 0.0 * IQR
filtered_Data = Data[(Data['Previous_Flight_Delay_Minutes'] >= lower_bound) &
    (Data['Previous_Flight_Delay_Minutes'] <= upper_bound)]
plt.figure(figsize=(8, 6))
sns.boxplot(x=filtered_Data['Previous_Flight_Delay_Minutes'])
plt.title('Boxplot of Previous_Flight_Delay_Minutes (Outliers Removed)')
plt.xlabel('Previous Flight Delay Minutes (Scaled)')
plt.show()
```




```
[26]: plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Airline_Rating'])
plt.title('Boxplot of Airline Rating')
plt.xlabel('Airline Rating (Scaled)')
plt.show()
```






```
[27]: plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Passenger_Load'])
plt.title('Boxplot of Passenger load')
plt.xlabel('Passenger load (Scaled)')
plt.show()
```



```
[28]: plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Flight_Cancelled'])
plt.title('Boxplot of Flight Cancelled')
plt.xlabel('Flight Cancelled (Scaled)')
plt.show()
```



```
[29]: data = pd.read_csv('Flyzy Flight Cancellation - Sheet1.csv')
```

```
[30]: data.describe()
```

```
[30]:
```

	Flight ID	Flight_Distance	Scheduled_Departure_Time	Day_of_Week \
count	3.000000e+03	3000.000000	3000.000000	3000.000000
mean	4.997429e+06	498.909333	11.435000	3.963000
std	2.868139e+06	98.892266	6.899298	2.016346
min	3.681000e+03	138.000000	0.000000	1.000000
25%	2.520313e+06	431.000000	6.000000	2.000000
50%	5.073096e+06	497.000000	12.000000	4.000000
75%	7.462026e+06	566.000000	17.000000	6.000000
max	9.999011e+06	864.000000	23.000000	7.000000

	Month	Weather_Score	Previous_Flight_Delay_Minutes \
count	3000.000000	3000.000000	3000.000000
mean	6.381000	0.524023	26.793383
std	3.473979	0.290694	27.874733

min	1.000000	0.000965	0.000000
25%	3.000000	0.278011	7.000000
50%	6.000000	0.522180	18.000000
75%	9.000000	0.776323	38.000000
max	12.000000	1.099246	259.000000

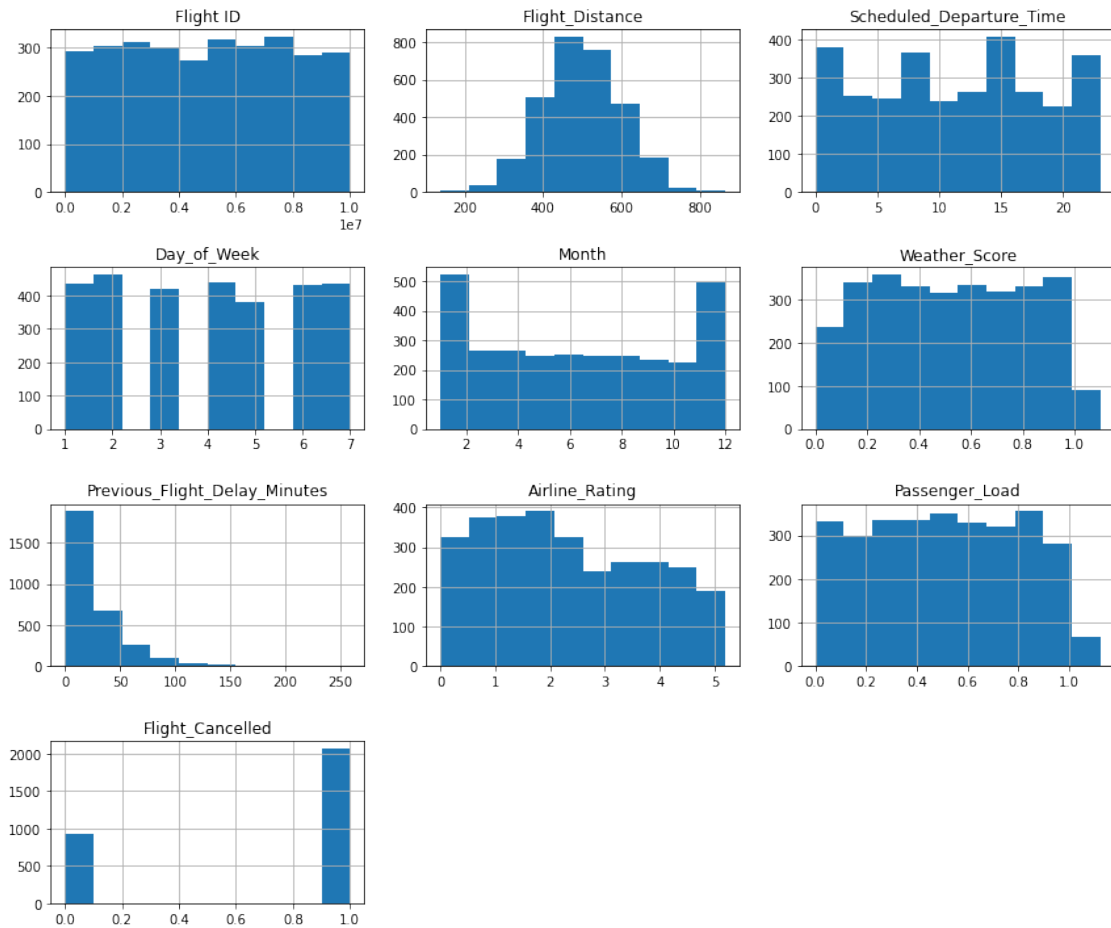
	Airline_Rating	Passenger_Load	Flight_Cancelled
count	3000.000000	3000.000000	3000.000000
mean	2.317439	0.515885	0.690667
std	1.430386	0.295634	0.462296
min	0.000103	0.001039	0.000000
25%	1.092902	0.265793	0.000000
50%	2.126614	0.517175	1.000000
75%	3.525746	0.770370	1.000000
max	5.189038	1.123559	1.000000

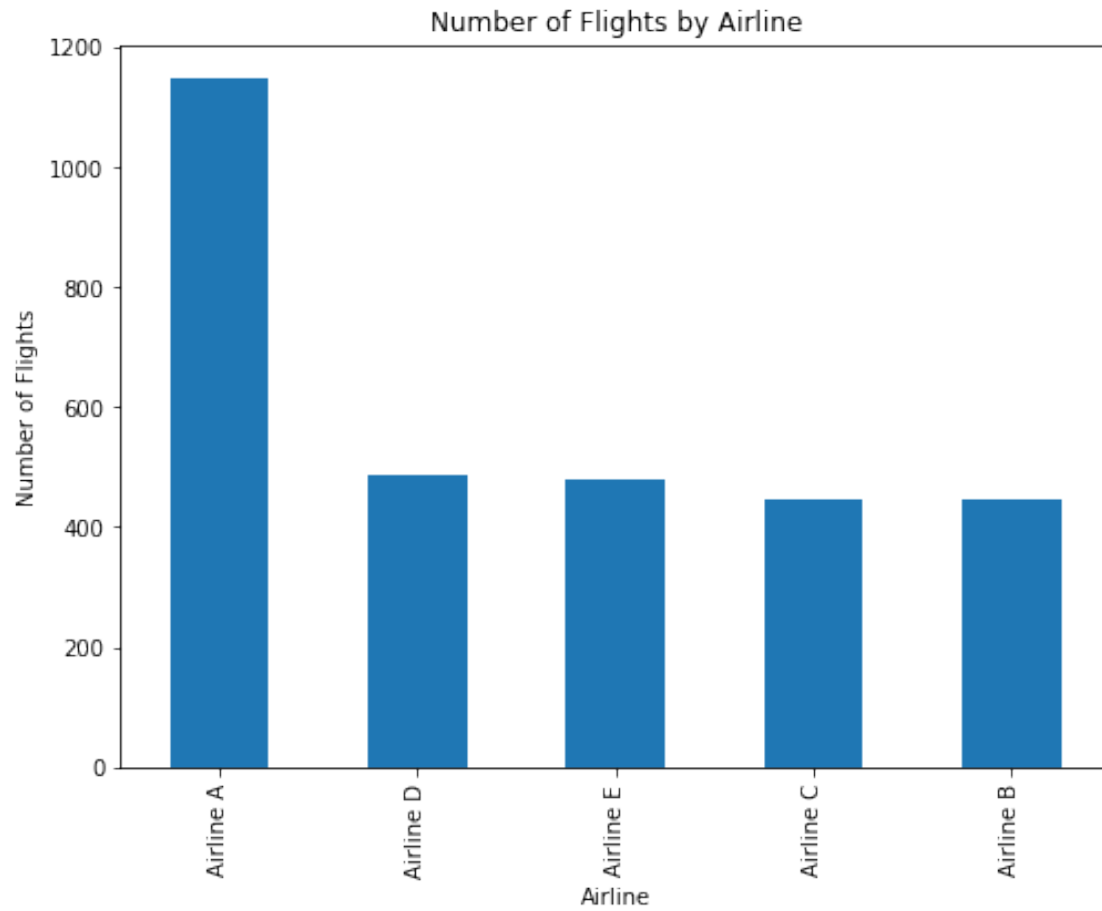
[31]: Data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Flight ID                            3000 non-null   int64
1   Airline                              3000 non-null   object
2   Flight_Distance                      3000 non-null   int64
3   Origin_Airport                      3000 non-null   object
4   Destination_Airport                 3000 non-null   object
5   Scheduled_Departure_Time             3000 non-null   int64
6   Day_of_Week                         3000 non-null   int64
7   Month                               3000 non-null   int64
8   Airplane_Type                       3000 non-null   object
9   Weather_Score                      3000 non-null   float64
10  Previous_Flight_Delay_Minutes       3000 non-null   float64
11  Airline_Rating                      3000 non-null   float64
12  Passenger_Load                      3000 non-null   float64
13  Flight_Cancelled                    3000 non-null   int64
dtypes: float64(4), int64(6), object(4)
memory usage: 328.2+ KB
```


```
[32]: Data.hist(figsize=(12, 10))
plt.tight_layout()
plt.show()
plt.figure(figsize=(8, 6))
Data['Airline'].value_counts().plot(kind='bar')
plt.title('Number of Flights by Airline')
plt.xlabel('Airline')
```

```
plt.ylabel('Number of Flights')
plt.show()
print("\n")
```






```
[33]: plt.figure(figsize=(8, 6))
Data['Airline'].value_counts().plot(kind='bar')
plt.title('Number of Flights by Airline')
plt.xlabel('Airline')
plt.ylabel('Number of Flights')
plt.show()
```




```
[34]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Flight_Distance'], kde=True, bins=20, color='blue')
plt.title('Histogram of Flight Distance')
plt.xlabel('Flight Distance')
plt.ylabel('Frequency')
plt.show()
```




```
[35]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Scheduled_Departure_Time'], kde=True, bins=20, color='blue')
plt.title('Histogram of Scheduled Departure Time')
plt.xlabel('Scheduled Departure Time')
plt.ylabel('Frequency')
plt.show()
```

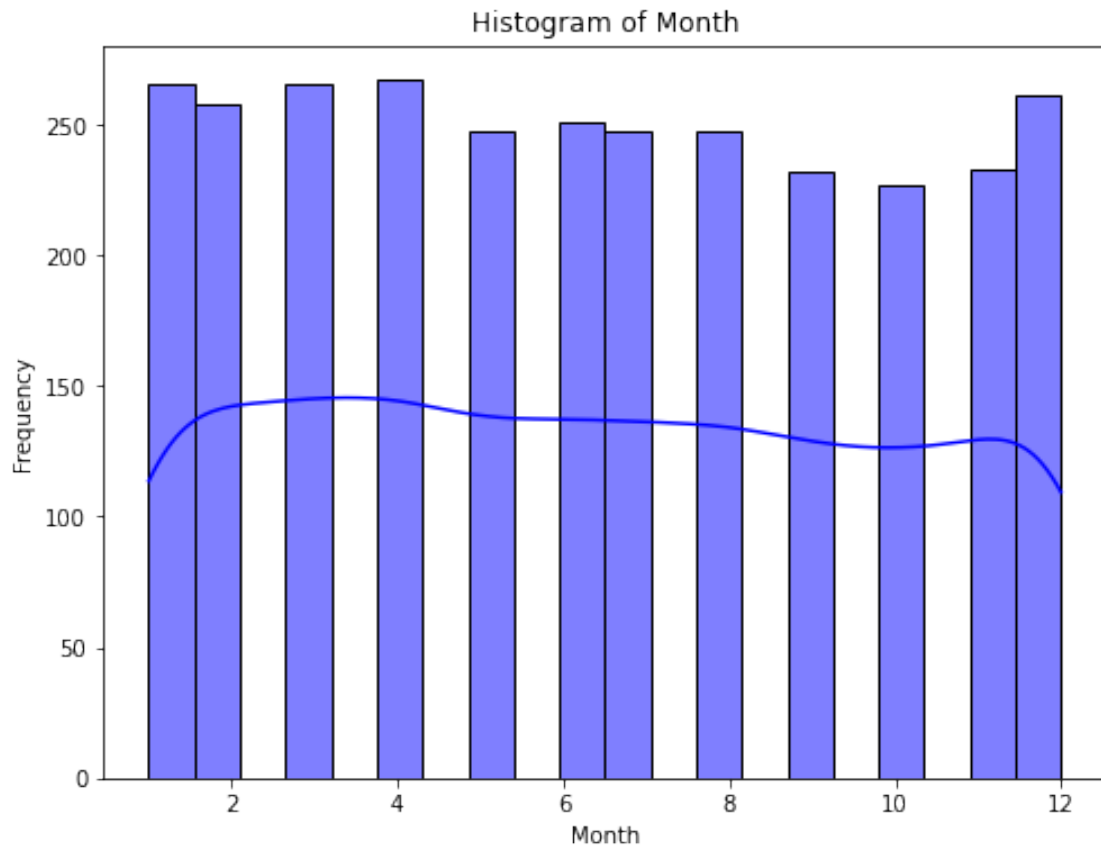





```
[36]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Day_of_Week'], kde=True, bins=20, color='blue')
plt.title('Histogram of Day of Week')
plt.xlabel('Day of Week')
plt.ylabel('Frequency')
plt.show()
```



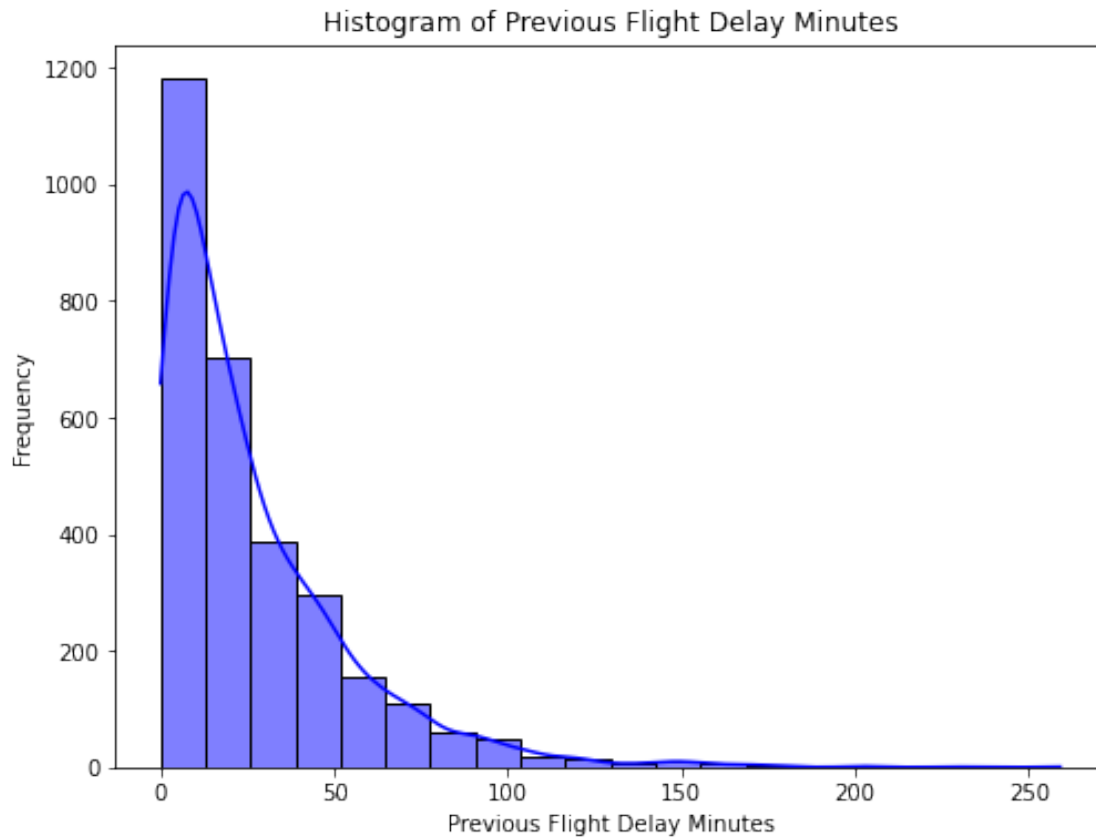
```
[37]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Month'], kde=True, bins=20, color='blue')
plt.title('Histogram of Month')
plt.xlabel('Month')
plt.ylabel('Frequency')
plt.show()
```




```
[38]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Weather_Score'], kde=True, bins=20, color='blue')
plt.title('Histogram of Weather_Score')
plt.xlabel('Weather_Score')
plt.ylabel('Frequency')
plt.show()
```




```
[39]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Previous_Flight_Delay_Minutes'], kde=True, bins=20, color='blue')
plt.title('Histogram of Previous Flight Delay Minutes')
plt.xlabel('Previous Flight Delay Minutes')
plt.ylabel('Frequency')
plt.show()
```



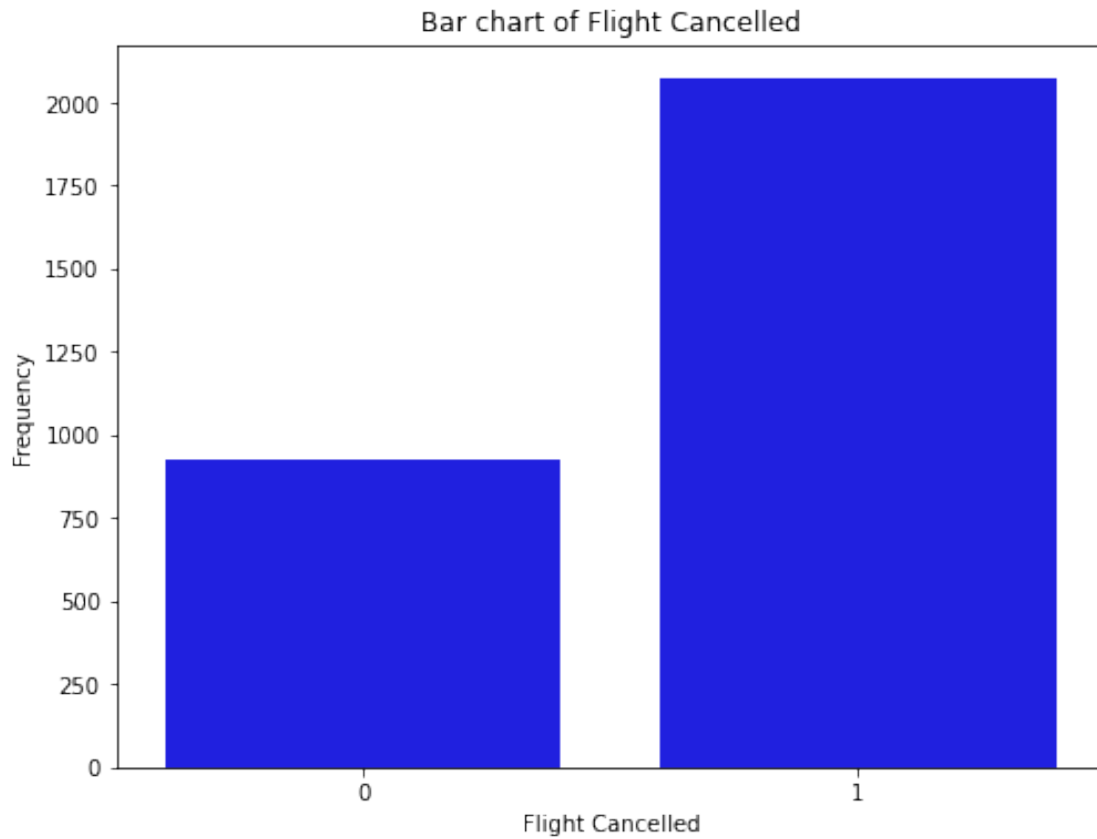
```
[40]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Airline_Rating'], kde=True, bins=20, color='blue')
plt.title('Histogram of Airline Rating')
plt.xlabel('Airline Rating')
plt.ylabel('Frequency')
plt.show()
```



```
[41]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Passenger_Load'], kde=True, bins=20, color='blue')
plt.title('Histogram of Passenger Load')
plt.xlabel('Passenger Load')
plt.ylabel('Frequency')
plt.show()
```

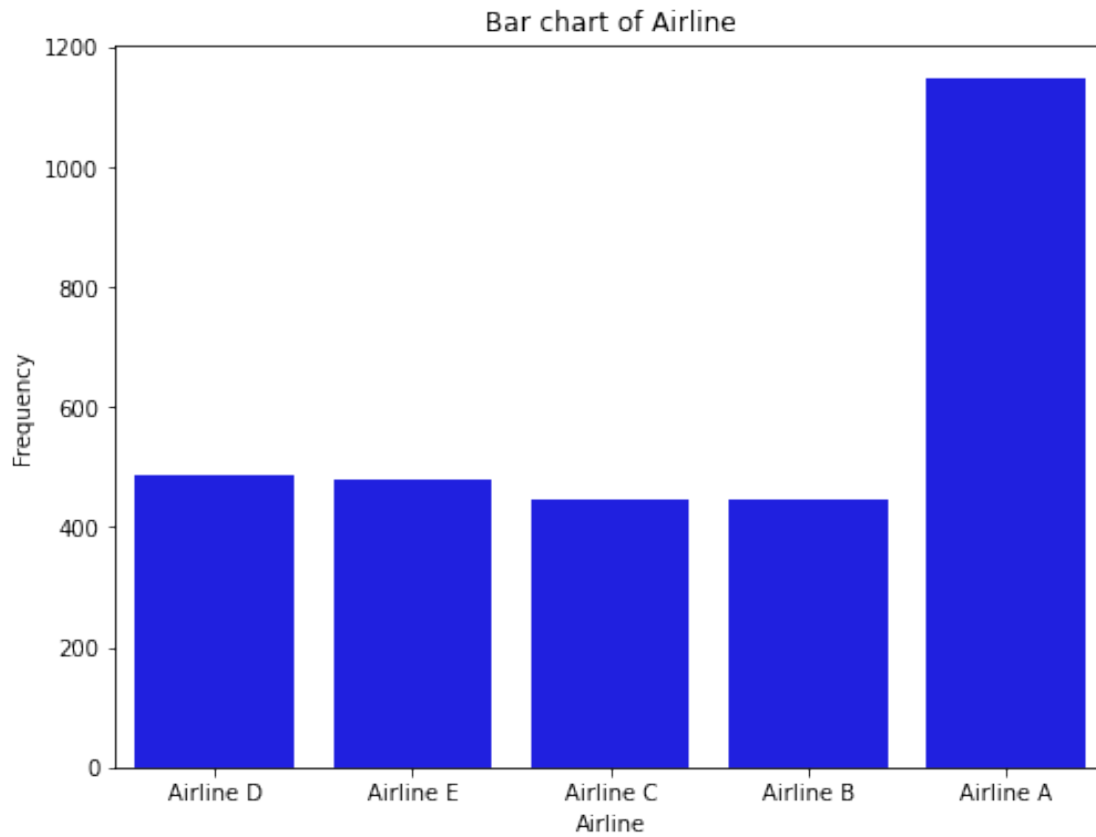


```
[42]: plt.figure(figsize=(8, 6))
sns.countplot(data=Data, x= 'Flight_Cancelled', color='blue')
plt.title('Bar chart of Flight Cancelled')
plt.xlabel('Flight Cancelled')
plt.ylabel('Frequency')
plt.show()
```




```
[43]: plt.figure(figsize=(8, 6))
sns.countplot(data=Data, x= 'Airline', color='blue')
plt.title('Bar chart of Airline')
plt.xlabel('Airline')
plt.ylabel('Frequency')
plt.show()
```

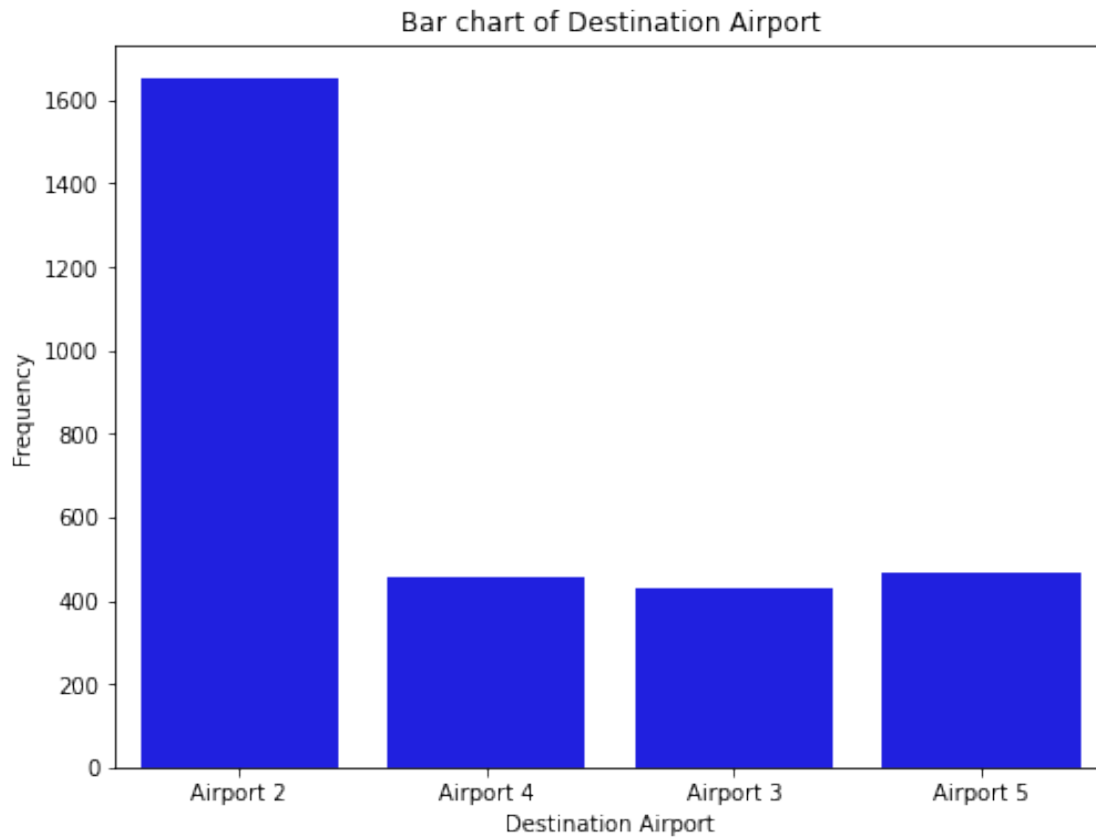




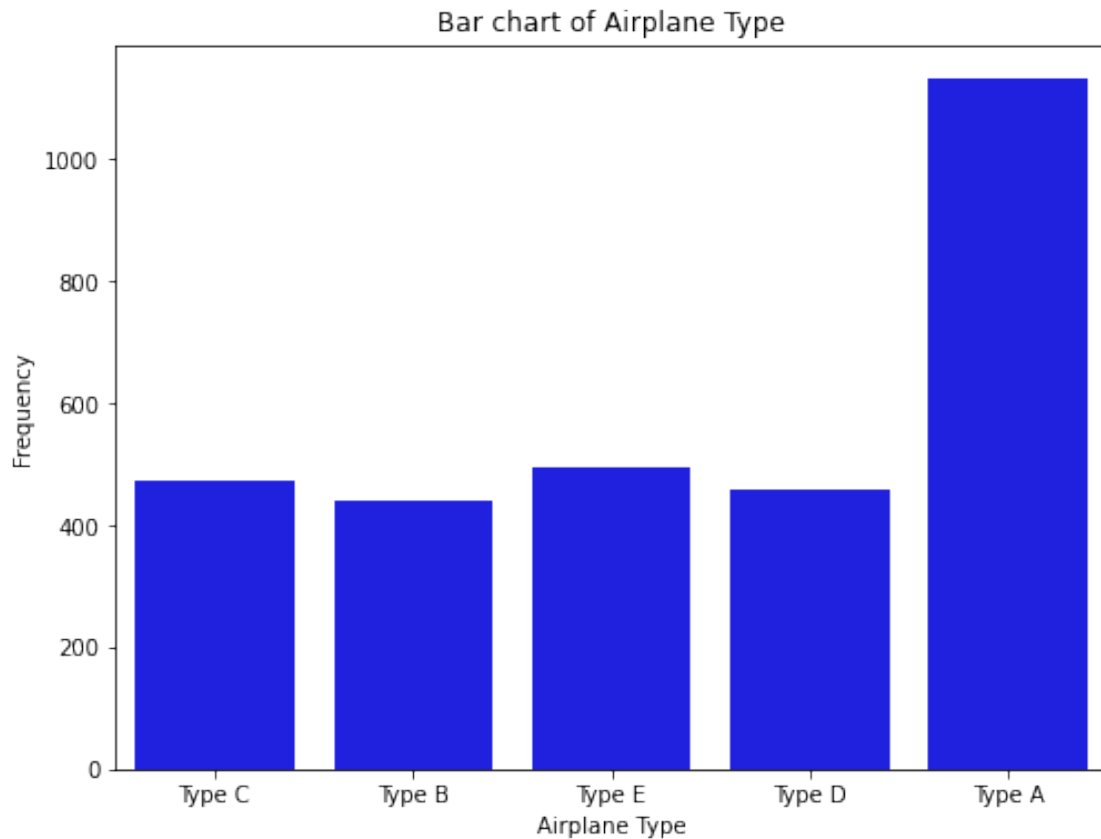
```
[44]: plt.figure(figsize=(8, 6))
sns.countplot(data=Data, x= 'Origin_Airport', color='blue')
plt.title('Bar chart of Origin Airport')
plt.xlabel('Origin Airport')
plt.ylabel('Frequency')
plt.show()
```



```
[45]: plt.figure(figsize=(8, 6))
sns.countplot(data=Data, x= 'Destination_Airport', color='blue')
plt.title('Bar chart of Destination Airport')
plt.xlabel('Destination Airport')
plt.ylabel('Frequency')
plt.show()
```



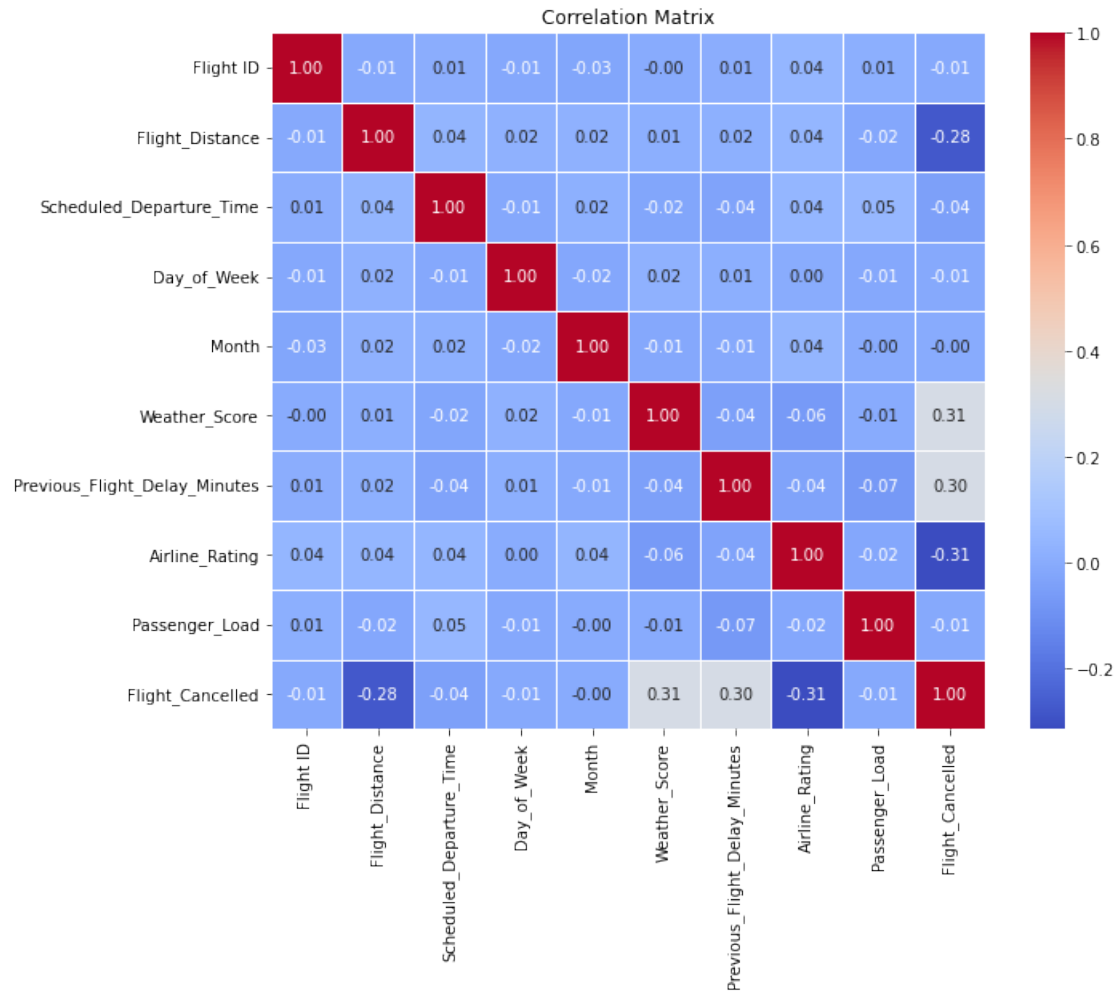
```
[46]: plt.figure(figsize=(8, 6))
sns.countplot(data=Data, x= 'Airplane_Type', color='blue')
plt.title('Bar chart of Airplane Type')
plt.xlabel('Airplane Type')
plt.ylabel('Frequency')
plt.show()
```



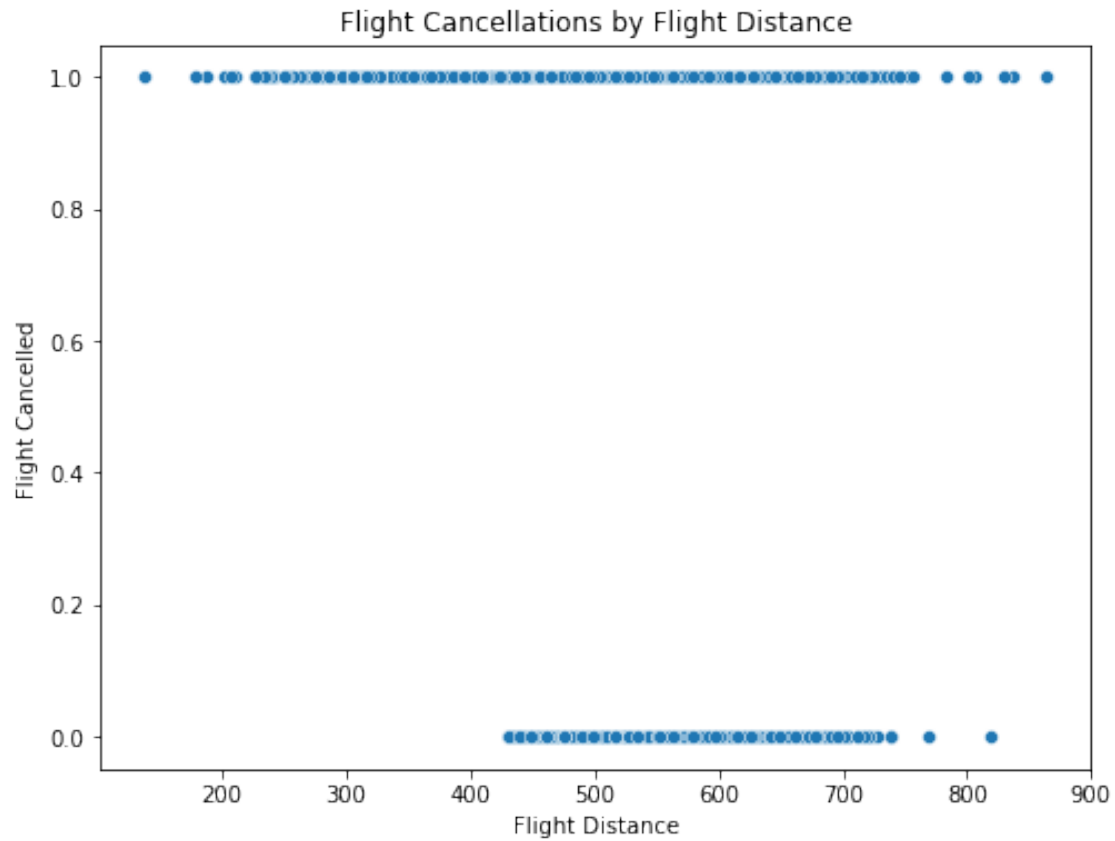
```
[47]: plt.figure(figsize=(10, 8))
sns.heatmap(Data.corr(), annot=True, cmap= 'coolwarm', fmt=".2f", linewidths=0.
↪5)
plt.title('Correlation Matrix')
plt.show()
print("\n")
```

/tmp/ipykernel\_91/3230157905.py:2: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.


```
sns.heatmap(Data.corr(), annot=True, cmap= 'coolwarm', fmt=".2f",
linewidths=0.5)
```




```
[48]: plt.figure(figsize=(8, 6))
sns.scatterplot(data=Data, x='Flight_Distance', y='Flight_Cancelled')
plt.title('Flight Cancellations by Flight Distance')
plt.xlabel('Flight Distance')
plt.ylabel('Flight Cancelled')
plt.show()
```




```
[49]: plt.figure(figsize=(12, 6))
sns.countplot(data=Data, x='Airline', hue='Flight_Cancelled')
plt.title('Flight Cancellations by Airline')
plt.xlabel('Airline')
plt.ylabel('Number of Flights')
plt.xticks(rotation=45)
plt.legend(title='Flight Cancelled', loc='upper right')
plt.show()
```



```
[50]: plt.figure(figsize=(12, 6))
sns.countplot(data=Data, x='Origin_Airport', hue='Flight_Cancelled')
plt.title('Flight Cancellations by Origin Airport')
plt.xlabel('Origin Airport')
plt.ylabel('Number of Flights')
plt.xticks(rotation=45)
plt.legend(title='Flight Cancelled', loc='upper right')
plt.show()
```



```
[51]: numeric_data = Data.select_dtypes(include='number')
correlation_with_target = numeric_data.corr()['Flight_Cancelled'].
    ↪sort_values(ascending=False)
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_with_target.to_frame(), annot=True, cmap='coolwarm',
    ↪fmt=".2f")
plt.title('Correlation of Numeric Features with Flight Cancelled')
plt.xlabel('Numeric Features')
plt.ylabel('Flight Cancelled')
plt.show()
```



## Preprocessing and Model Building ##

```
[52]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, MinMaxScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↪f1_score, roc_auc_score
```



```
[53]: X = Data.drop(['Flight ID', 'Flight_Cancelled'], axis=1) # Features
      y = Data['Flight_Cancelled']
```

```
[59]: x_encoded = pd.get_dummies(X)
```

```
[61]: print(x_encoded.head())
```

	Flight_Distance	Scheduled_Departure_Time	Day_of_Week	Month	\
0	475		4	6	1
1	538		12	1	6
2	565		17	3	9
3	658		1	1	8
4	566		19	7	12

	Weather_Score	Previous_Flight_Delay_Minutes	Airline_Rating	\
0	0.225122		5.0	2.151974
1	0.060346		68.0	1.600779
2	0.093920		18.0	4.406848
3	0.656750		13.0	0.998757
4	0.505211		4.0	3.806206

	Passenger_Load	Airline_Airline A	Airline_Airline B	...	\
0	0.477202	0	0	...	
1	0.159718	0	0	...	
2	0.256803	0	0	...	
3	0.504077	0	0	...	
4	0.019638	0	0	...	

	Origin_Airport_Airport 5	Destination_Airport_Airport 2	\
0	0	1	
1	1	0	
2	0	1	
3	1	0	
4	0	1	

	Destination_Airport_Airport 3	Destination_Airport_Airport 4	\
0	0	0	
1	0	1	
2	0	0	
3	1	0	
4	0	0	

	Destination_Airport_Airport 5	Airplane_Type_Type A	Airplane_Type_Type B	\
0	0	0	0	
1	0	0	1	
2	0	0	0	
3	0	0	1	

4 0 0 0

	Airplane_Type_Type C	Airplane_Type_Type D	Airplane_Type_Type E
0	1	0	0
1	0	0	0
2	1	0	0
3	0	0	0
4	0	0	1

[5 rows x 27 columns]

```
[63]: scaler = MinMaxScaler()
      x_scaled = scaler.fit_transform(x_encoded)
```

```
[64]: print(x_scaled)
```

```
[[0.46418733 0.17391304 0.83333333 ... 1.      0.      0.      ]
 [0.55096419 0.52173913 0.          ... 0.      0.      0.      ]
 [0.58815427 0.73913043 0.33333333 ... 1.      0.      0.      ]
 ...
 [0.44490358 0.34782609 0.33333333 ... 0.      0.      0.      ]
 [0.44903581 0.2173913  0.66666667 ... 0.      0.      1.      ]
 [0.31818182 0.04347826 0.          ... 0.      0.      0.      ]]
```

```
[66]: from sklearn.model_selection import train_test_split
      x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size=0.2,
      random_state=42)
      print("Training set - Features:", x_train.shape, "Target:", y_train.shape)
      print("Test set - Features:", x_test.shape, "Target:", y_test.shape)
```

```
Training set - Features: (2400, 27) Target: (2400,)
Test set - Features: (600, 27) Target: (600,)
```

```
[67]: from sklearn.linear_model import LogisticRegression
```

```
[68]: logreg_model = LogisticRegression(random_state=42)
```

```
[70]: logreg_model.fit(x_train, y_train)
```

```
[70]: LogisticRegression(random_state=42)
```

```
[71]: print("Model Coefficients:", logreg_model.coef_)
```

```
Model Coefficients: [[-5.44813124  0.01029947 -0.08940905  0.16600798
 3.36691911  8.39711858
 -2.91703632  0.05163797 -0.02845402 -0.14100312  0.0460956  0.03333066
 0.09012382 -0.02711148 -0.05280328  0.01537221  0.10468177 -0.04004627
```

```
0.05851232 -0.03401603 0.07954138 -0.10394473 0.04275748 -0.21807963
0.02008184 0.22047041 -0.06513717]]
```

```
[72]: from sklearn.metrics import accuracy_score, precision_score, recall_score, \
      ↪ f1_score, roc_auc_score
```

```
[73]: y_pred = logreg_model.predict(x_test)
```

```
[74]: from sklearn.metrics import accuracy_score, precision_score, recall_score, \
      ↪ f1_score, roc_auc_score
```

## 0.2 Build other Classification Models

```
[75]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler, OneHotEncoder
      from sklearn.compose import ColumnTransformer
      from sklearn.pipeline import Pipeline
```

```
[76]: X = Data.drop(columns=['Scheduled_Departure_Time', 'Flight_Cancelled'])
      y = Data['Flight_Cancelled']
      numerical_features = ['Flight_Distance', 'Scheduled_Departure_Time', \
      ↪ 'Previous_Flight_Delay_Minutes', 'Airline_Rating', 'Passenger_Load', \
      ↪ 'Weather_Score']
      categorical_features = ['Airline', 'Origin_Airport', 'Destination_Airport', \
      ↪ 'Airplane_Type', 'Day_of_Week', 'Month']
      preprocessor = ColumnTransformer(
          transformers=[
              ('num', StandardScaler(), numerical_features),
              ('cat', OneHotEncoder(), categorical_features)
          ])
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, \
      ↪ random_state=42)
```

```
[77]: numerical_features = ['Flight_Distance', 'Sched_Departure_Time', \
      ↪ 'Previous_Flight_Delay_Minutes', 'Airline_Rating', 'Passenger_Load', \
      ↪ 'Weather_Score']
```

```
[80]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler, OneHotEncoder
      from sklearn.compose import ColumnTransformer
      from sklearn.pipeline import Pipeline
      from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import classification_report, roc_auc_score, \
      ↪ precision_score, recall_score, f1_score
```

```
[ ]:
```

```
[81]: data = pd.read_csv('Flyzy Flight Cancellation - Sheet1.csv')
data.iloc[1:4, 2:3]= np.NaN
data.iloc[1:4, 3:4]= "NA"
data.iloc[1:4, 4:5]= ""
data["None_col"]= None
data.head()
```

```
[81]:
```

	Flight ID	Airline	Flight_Distance	Origin_Airport	Destination_Airport	\
0	7319483	Airline D	475.0	Airport 3	Airport 2	
1	4791965	Airline E	NaN	NA		
2	2991718	Airline C	NaN	NA		
3	4220106	Airline E	NaN	NA		
4	2263008	Airline E	566.0	Airport 2	Airport 2	

	Scheduled_Departure_Time	Day_of_Week	Month	Airplane_Type	Weather_Score	\
0		4	6	1	Type C	0.225122
1		12	1	6	Type B	0.060346
2		17	3	9	Type C	0.093920
3		1	1	8	Type B	0.656750
4		19	7	12	Type E	0.505211

	Previous_Flight_Delay_Minutes	Airline_Rating	Passenger_Load	\
0		5.0	2.151974	0.477202
1		68.0	1.600779	0.159718
2		18.0	4.406848	0.256803
3		13.0	0.998757	0.504077
4		4.0	3.806206	0.019638

	Flight_Cancelled	None_col
0	0	None
1	1	None
2	0	None
3	1	None
4	0	None

```
[83]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, roc_auc_score, \
    precision_score, recall_score, f1_score
```

```

Data = pd.read_csv('Flyzy Flight Cancellation - Sheet1.csv')

print(Data.columns)

X = Data.drop(columns=['Flight_Cancelled', 'Flight_Cancelled'])
y = Data['Flight_Cancelled']

numerical_features = ['Flight_Distance', 'Scheduled_Departure_Time',
    ↪ 'Previous_Flight_Delay_Minutes', 'Airline_Rating', 'Passenger_Load',
    ↪ 'Weather_Score']
categorical_features = ['Airline', 'Origin_Airport', 'Destination_Airport',
    ↪ 'Airplane_Type', 'Day_of_Week', 'Month']

preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numerical_features),
        ('cat', OneHotEncoder(), categorical_features)
    ])

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=42)

```

```

Index(['Flight ID', 'Airline', 'Flight_Distance', 'Origin_Airport',
      'Destination_Airport', 'Scheduled_Departure_Time', 'Day_of_Week',
      'Month', 'Airplane_Type', 'Weather_Score',
      'Previous_Flight_Delay_Minutes', 'Airline_Rating', 'Passenger_Load',
      'Flight_Cancelled'],
      dtype='object')

```

```

[84]: numerical_features = ['Flight_Distance', 'Sched_Departure_Time',
    ↪ 'Previous_Flight_Delay_Minutes', 'Airline_Rating', 'Passenger_Load',
    ↪ 'Weather_Score']

```

```

[85]: lr_pipeline = Pipeline(steps=[('preprocessor', preprocessor),
    ('classifier',
    ↪ LogisticRegression(random_state=42))])

lr_pipeline.fit(X_train, y_train)

```

```

[85]: Pipeline(steps=[('preprocessor',
    ColumnTransformer(transformers=[('num', StandardScaler(),
    ['Flight_Distance',

```

```

'Scheduled_Departure_Time',
'Previous_Flight_Delay_Minutes',
'Airline_Rating',
'Passenger_Load',
'Weather_Score']],
('cat', OneHotEncoder(),
 ['Airline', 'Origin_Airport',
 'Destination_Airport',
 'Airplane_Type',
 'Day_of_Week', 'Month'])),
('classifier', LogisticRegression(random_state=42)))

```

```

[86]: from sklearn.linear_model import LogisticRegression

# Build pipeline
lr_pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                              ('classifier', LogisticRegression())])

# Train and evaluate
lr_pipeline.fit(X_train, y_train)
y_pred = lr_pipeline.predict(X_test)

# Evaluation
print("Logistic Regression Classification Report:")
print(classification_report(y_test, y_pred))
print("ROC-AUC Score:", roc_auc_score(y_test, lr_pipeline.
    ↪predict_proba(X_test)[:, 1]))

```

```

Logistic Regression Classification Report:

```

	precision	recall	f1-score	support
0	0.70	0.59	0.64	187
1	0.83	0.89	0.86	413
accuracy			0.79	600
macro avg	0.76	0.74	0.75	600
weighted avg	0.79	0.79	0.79	600

```

ROC-AUC Score: 0.8652613587808006

```

```

[89]: from sklearn.tree import DecisionTreeClassifier

dt_pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                              ('classifier',
    ↪DecisionTreeClassifier(random_state=42))])

```

```

dt_pipeline.fit(X_train, y_train)
y_pred = dt_pipeline.predict(X_test)

print("Decision Tree Classification Report:")
print(classification_report(y_test, y_pred))
print("ROC-AUC Score:", roc_auc_score(y_test, dt_pipeline.
    ↳predict_proba(X_test)[: , 1]))

```

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	0.97	0.94	0.95	187
1	0.97	0.99	0.98	413
accuracy			0.97	600
macro avg	0.97	0.96	0.96	600
weighted avg	0.97	0.97	0.97	600

ROC-AUC Score: 0.9606505159845139

```

[90]: from sklearn.ensemble import GradientBoostingClassifier

gb_pipeline = Pipeline(steps=[('preprocessor', preprocessor),
    ('classifier',
    ↳GradientBoostingClassifier(random_state=42))])

gb_pipeline.fit(X_train, y_train)
y_pred = gb_pipeline.predict(X_test)

print("Gradient Boosting Classification Report:")
print(classification_report(y_test, y_pred))
print("ROC-AUC Score:", roc_auc_score(y_test, gb_pipeline.
    ↳predict_proba(X_test)[: , 1]))

```

Gradient Boosting Classification Report:

	precision	recall	f1-score	support
0	0.96	1.00	0.98	187
1	1.00	0.98	0.99	413
accuracy			0.99	600
macro avg	0.98	0.99	0.98	600
weighted avg	0.99	0.99	0.99	600

ROC-AUC Score: 0.9917779130142041

```
[98]: from sklearn.ensemble import GradientBoostingClassifier

gb_pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                               ('classifier',
                                ↪GradientBoostingClassifier(random_state=42))])

gb_pipeline.fit(X_train, y_train)
y_pred = gb_pipeline.predict(X_test)

print("Gradient Boosting Classification Report:")
print(classification_report(y_test, y_pred))
print("ROC-AUC Score:", roc_auc_score(y_test, gb_pipeline.
    ↪predict_proba(X_test)[: , 1]))
```

Gradient Boosting Classification Report:

	precision	recall	f1-score	support
0	0.96	1.00	0.98	187
1	1.00	0.98	0.99	413
accuracy			0.99	600
macro avg	0.98	0.99	0.98	600
weighted avg	0.99	0.99	0.99	600

ROC-AUC Score: 0.9917779130142041

```
[1]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.metrics import precision_score, recall_score, f1_score,
    ↪roc_auc_score, classification_report
import matplotlib.pyplot as plt

Data = pd.read_csv('Flyzy Flight Cancellation - Sheet1.csv')
```



```

X = Data.drop(columns=['Flight_Cancelled', 'Flight_Cancelled'])
y = Data['Flight_Cancelled']

numerical_features = ['Flight_Distance', 'Scheduled_Departure_Time',
    ↪ 'Previous_Flight_Delay_Minutes', 'Airline_Rating', 'Passenger_Load',
    ↪ 'Weather_Score']
categorical_features = ['Airline', 'Origin_Airport', 'Destination_Airport',
    ↪ 'Airplane_Type', 'Day_of_Week', 'Month']

preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numerical_features),
        ('cat', OneHotEncoder(), categorical_features)
    ])

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=42)

models = {
    "Logistic Regression": LogisticRegression(random_state=42),
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "Random Forest": RandomForestClassifier(random_state=42),
    "Gradient Boosting": GradientBoostingClassifier(random_state=42),
    "SVM": SVC(probability=True, random_state=42)
}

results = {}

for model_name, model in models.items():
    pipeline = Pipeline(steps=[('preprocessor', preprocessor),
        ↪ ('classifier', model)])
    pipeline.fit(X_train, y_train)
    y_pred = pipeline.predict(X_test)
    y_proba = pipeline.predict_proba(X_test)[:, 1] if hasattr(model,
    ↪ "predict_proba") else None
    results[model_name] = {
        "Precision": precision_score(y_test, y_pred),
        "Recall": recall_score(y_test, y_pred),
        "F1 Score": f1_score(y_test, y_pred),
        "ROC-AUC": roc_auc_score(y_test, y_proba) if y_proba is not None else
    ↪ None


```

```
}
```

```
results_df = pd.DataFrame(results).T  
print(results_df)
```

	Precision	Recall	F1 Score	ROC-AUC
Logistic Regression	0.826185	0.886199	0.855140	0.865261
Decision Tree	0.971360	0.985472	0.978365	0.960651
Random Forest	0.987805	0.980630	0.984204	0.990703
Gradient Boosting	1.000000	0.980630	0.990220	0.991778
SVM	0.913551	0.946731	0.929845	0.970996

```
[2]: results_df.plot(kind='bar', figsize=(12, 8))  
plt.title('Model Comparison')  
plt.ylabel('Score')  
plt.xticks(rotation=45)  
plt.show()
```



```
[ ]:
```