

Preprocessing and Model Building

September 16, 2024

```
[1]: import pandas as pd
import numpy as np
```

```
[2]: data = pd.read_csv('Flyzy Flight Cancellation - Sheet1.csv')
data.iloc[1:4, 2:3]= np.NaN
data.iloc[1:4, 3:4]= "NA"
data.iloc[1:4, 4:5]= ""
data["None_col"]= None
data.head()
```

```
[2]:   Flight ID   Airline Flight_Distance Origin_Airport Destination_Airport \
0    7319483  Airline D           475.0      Airport 3      Airport 2
1    4791965  Airline E              NaN              NA
2    2991718  Airline C              NaN              NA
3    4220106  Airline E              NaN              NA
4    2263008  Airline E           566.0      Airport 2      Airport 2
```

```
   Scheduled_Departure_Time  Day_of_Week  Month Airplane_Type  Weather_Score \
0                        4              6      1      Type C      0.225122
1                       12              1      6      Type B      0.060346
2                       17              3      9      Type C      0.093920
3                        1              1      8      Type B      0.656750
4                       19              7     12      Type E      0.505211
```

```
   Previous_Flight_Delay_Minutes  Airline_Rating  Passenger_Load \
0                        5.0        2.151974        0.477202
1                       68.0        1.600779        0.159718
2                       18.0        4.406848        0.256803
3                       13.0        0.998757        0.504077
4                        4.0        3.806206        0.019638
```

```
   Flight_Cancelled  None_col
0                0      None
1                1      None
2                0      None
3                1      None
4                0      None
```

```
[3]: null= pd.isnull(data)
      null.head()
```

```
[3]: Flight ID  Airline  Flight_Distance  Origin_Airport  Destination_Airport  \
0      False    False                False            False            False
1      False    False                True             False            False
2      False    False                True             False            False
3      False    False                True             False            False
4      False    False                False            False            False

      Scheduled_Departure_Time  Day_of_Week  Month  Airplane_Type  Weather_Score  \
0                          False         False  False          False          False
1                          False         False  False          False          False
2                          False         False  False          False          False
3                          False         False  False          False          False
4                          False         False  False          False          False

      Previous_Flight_Delay_Minutes  Airline_Rating  Passenger_Load  \
0                          False         False          False
1                          False         False          False
2                          False         False          False
3                          False         False          False
4                          False         False          False

      Flight_Cancelled  None_col
0              False      True
1              False      True
2              False      True
3              False      True
4              False      True
```

```
[4]: data.fillna(0).head()
```

```
[4]: Flight ID  Airline  Flight_Distance  Origin_Airport  Destination_Airport  \
0      7319483  Airline D          475.0          Airport 3          Airport 2
1      4791965  Airline E           0.0              NA
2      2991718  Airline C           0.0              NA
3      4220106  Airline E           0.0              NA
4      2263008  Airline E          566.0          Airport 2          Airport 2

      Scheduled_Departure_Time  Day_of_Week  Month  Airplane_Type  Weather_Score  \
0                          4              6      1          Type C      0.225122
1                         12              1      6          Type B      0.060346
2                         17              3      9          Type C      0.093920
3                          1              1      8          Type B      0.656750
4                         19              7     12          Type E      0.505211
```

	Previous_Flight_Delay_Minutes	Airline_Rating	Passenger_Load \
0	5.0	2.151974	0.477202
1	68.0	1.600779	0.159718
2	18.0	4.406848	0.256803
3	13.0	0.998757	0.504077
4	4.0	3.806206	0.019638

	Flight_Cancelled	None_col
0	0	0
1	1	0
2	0	0
3	1	0
4	0	0


```
[5]: pd.isnull(data).sum().sum()
```

```
[5]: 3003
```

```
[7]: missing_values = data.isnull().sum()
print("Missing values per column (before handling):")
print(missing_values)
```

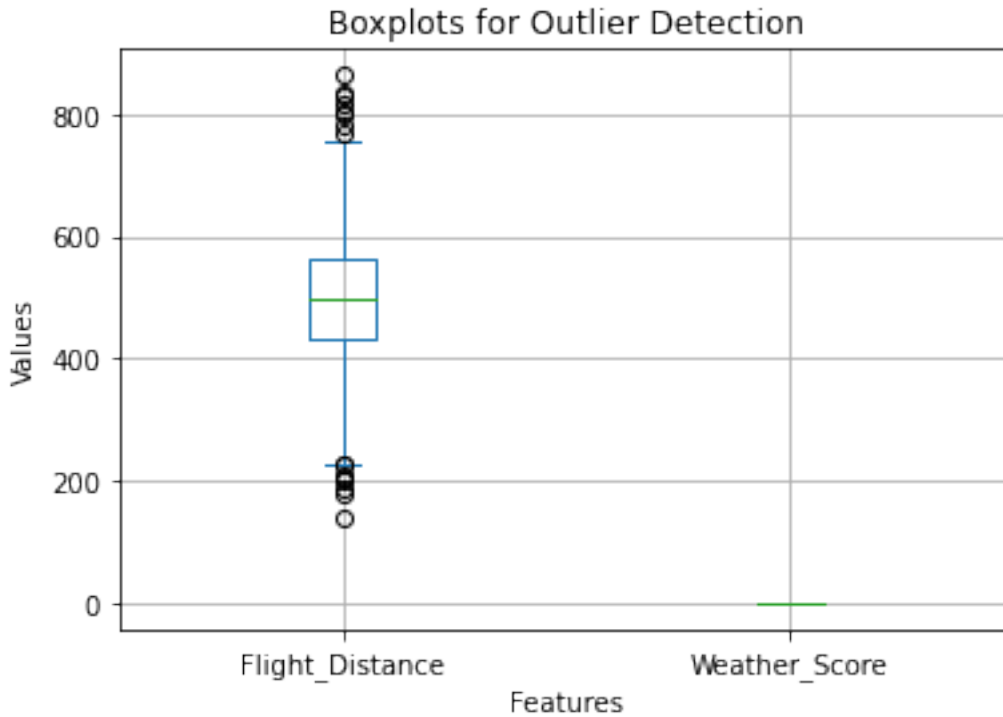
```
Missing values per column (before handling):
Flight ID                0
Airline                  0
Flight_Distance          3
Origin_Airport           0
Destination_Airport      0
Scheduled_Departure_Time 0
Day_of_Week              0
Month                    0
Airplane_Type            0
Weather_Score            0
Previous_Flight_Delay_Minutes 0
Airline_Rating           0
Passenger_Load          0
Flight_Cancelled         0
None_col                 3000
dtype: int64
```

```
[8]: import pandas as pd
import matplotlib.pyplot as plt
data.plot(kind='scatter' , x= 'Flight_Distance', y= 'Weather_Score')
plt.show()
```



```
[9]: import matplotlib.pyplot as plt
import pandas as pd
columns_to_check = ['Flight_Distance', 'Weather_Score']
plt.figure(figsize=(10, 6))
data[columns_to_check].plot(kind= 'box')
plt.title('Boxplots for Outlier Detection')
plt.xlabel('Features')
plt.ylabel('Values')
plt.grid(True)
plt.show()
```

<Figure size 720x432 with 0 Axes>



```
[10]: data_types = data.dtypes
      print("Data types of each column:")
      print(data_types)
```

```
Data types of each column:
Flight ID          int64
Airline            object
Flight_Distance    float64
Origin_Airport     object
Destination_Airport object
Scheduled_Departure_Time int64
Day_of_Week        int64
Month              int64
Airplane_Type      object
Weather_Score      float64
Previous_Flight_Delay_Minutes float64
Airline_Rating     float64
Passenger_Load     float64
Flight_Cancelled   int64
None_col           object
dtype: object
```

```
[11]: Data =pd.read_csv('Flyzy Flight Cancellation - Sheet1.csv')
```

```
[12]: print(Data.head())
```

	Flight ID	Airline	Flight_Distance	Origin_Airport	Destination_Airport	\
0	7319483	Airline D	475	Airport 3	Airport 2	
1	4791965	Airline E	538	Airport 5	Airport 4	
2	2991718	Airline C	565	Airport 1	Airport 2	
3	4220106	Airline E	658	Airport 5	Airport 3	
4	2263008	Airline E	566	Airport 2	Airport 2	

	Scheduled_Departure_Time	Day_of_Week	Month	Airplane_Type	Weather_Score	\
0		4	6	1	Type C	0.225122
1		12	1	6	Type B	0.060346
2		17	3	9	Type C	0.093920
3		1	1	8	Type B	0.656750
4		19	7	12	Type E	0.505211

	Previous_Flight_Delay_Minutes	Airline_Rating	Passenger_Load	\
0	5.0	2.151974	0.477202	
1	68.0	1.600779	0.159718	
2	18.0	4.406848	0.256803	
3	13.0	0.998757	0.504077	
4	4.0	3.806206	0.019638	

	Flight_Cancelled
0	0
1	1
2	0
3	1
4	0

```
[13]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[14]: Data =pd.read_csv('Flyzy Flight Cancellation - Sheet1.csv')
```

```
[15]: print(Data.head())
```

	Flight ID	Airline	Flight_Distance	Origin_Airport	Destination_Airport	\
0	7319483	Airline D	475	Airport 3	Airport 2	
1	4791965	Airline E	538	Airport 5	Airport 4	
2	2991718	Airline C	565	Airport 1	Airport 2	
3	4220106	Airline E	658	Airport 5	Airport 3	
4	2263008	Airline E	566	Airport 2	Airport 2	

	Scheduled_Departure_Time	Day_of_Week	Month	Airplane_Type	Weather_Score	\
0		4	6	1	Type C	0.225122

1	12	1	6	Type B	0.060346
2	17	3	9	Type C	0.093920
3	1	1	8	Type B	0.656750
4	19	7	12	Type E	0.505211

	Previous_Flight_Delay_Minutes	Airline_Rating	Passenger_Load \
0	5.0	2.151974	0.477202
1	68.0	1.600779	0.159718
2	18.0	4.406848	0.256803
3	13.0	0.998757	0.504077
4	4.0	3.806206	0.019638

	Flight_Cancelled
0	0
1	1
2	0
3	1
4	0

```
[16]: print(data.isnull().sum())
```

```
Flight ID          0
Airline            0
Flight_Distance    3
Origin_Airport     0
Destination_Airport 0
Scheduled_Departure_Time 0
Day_of_Week        0
Month              0
Airplane_Type      0
Weather_Score      0
Previous_Flight_Delay_Minutes 0
Airline_Rating     0
Passenger_Load     0
Flight_Cancelled   0
None_col           3000
dtype: int64
```

```
[17]: print(Data.dtypes)
```

```
Flight ID          int64
Airline            object
Flight_Distance    int64
Origin_Airport     object
Destination_Airport object
Scheduled_Departure_Time int64
Day_of_Week        int64
Month              int64
```

```


Airplane_Type          object
Weather_Score          float64
Previous_Flight_Delay_Minutes float64
Airline_Rating         float64
Passenger_Load         float64
Flight_Cancelled       int64
dtype: object

```

```

[18]: plt.figure(figsize=(12,8))
      sns.boxplot(data=Data)
      plt.xticks(rotation=45)
      plt.show()

```




```

[19]: import seaborn as sns
      import matplotlib.pyplot as plt
      plt.figure(figsize=(8, 6))
      sns.boxplot(x=Data['Flight_Distance'])
      plt.title('Boxplot of Flight Distance')


```




```
plt.xlabel('Flight Distance (Scaled)')
plt.show()
```




```
[20]: Q1 = Data['Flight_Distance'].quantile(0.25)
      Q3 = Data['Flight_Distance'].quantile(0.75)
      IQR = Q3 - Q1
      lower_bound = Q1 - 1.5 * IQR
      upper_bound = Q3 + 1.5 * IQR
      filtered_Data = Data[(Data['Flight_Distance'] >= lower_bound) &
                           (Data['Flight_Distance'] <= upper_bound)]
      plt.figure(figsize=(8, 6))
      sns.boxplot(x=filtered_Data['Flight_Distance'])
      plt.title('Boxplot of flight Distance (Outliers Removed)')
      plt.xlabel('Flight Distance (Scaled)')
      plt.show()
```



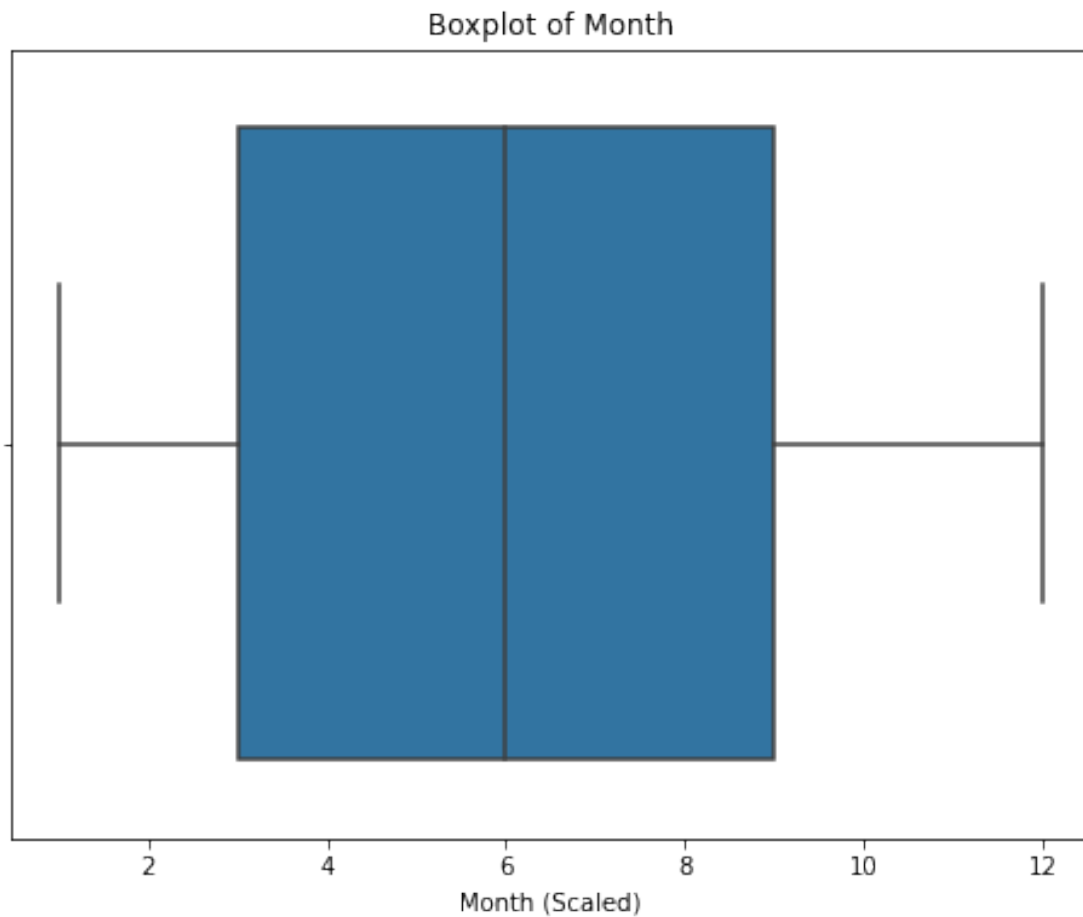
```
[21]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Scheduled_Departure_Time'])
plt.title('Boxplot of Scheduled Departure Time')
plt.xlabel('Scheduled Departure Time (Scaled)')
plt.show()
```



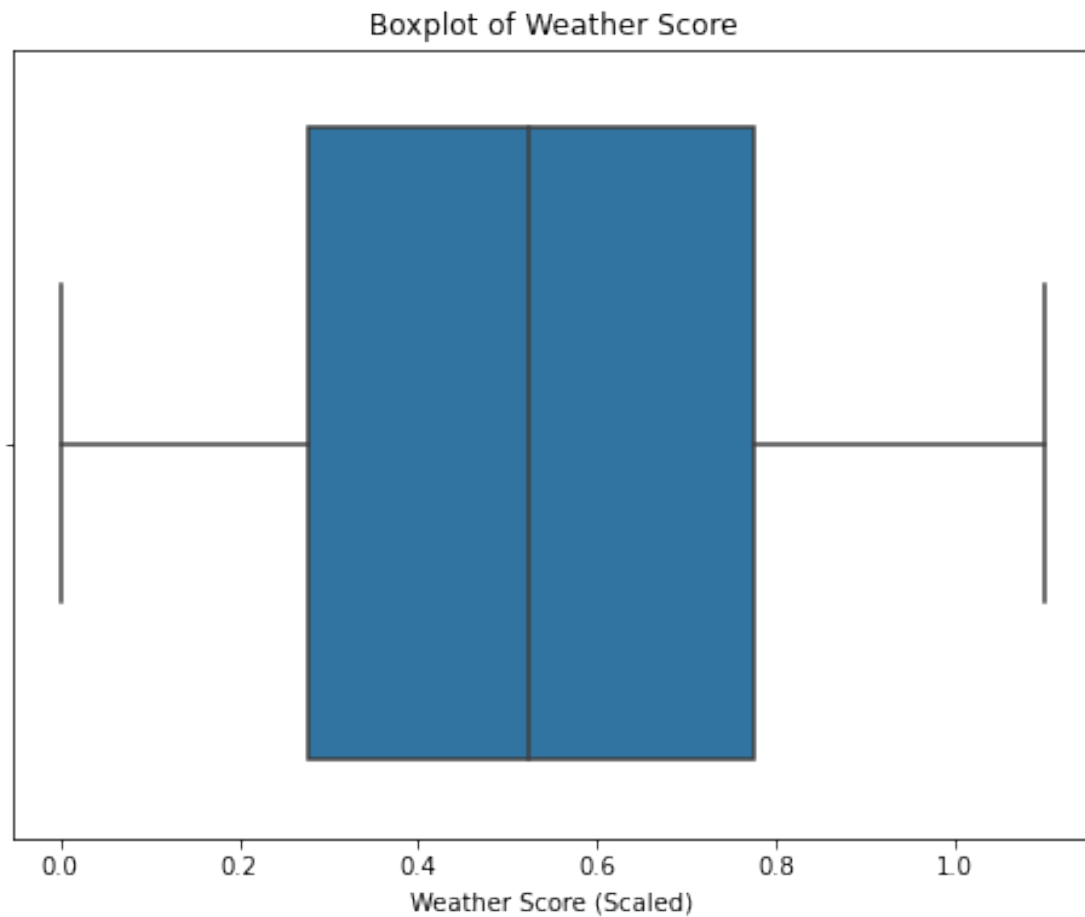
```
[22]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Day_of_Week'])
plt.title('Boxplot of Day of Week')
plt.xlabel('Day of Week (Scaled)')
plt.show()
```



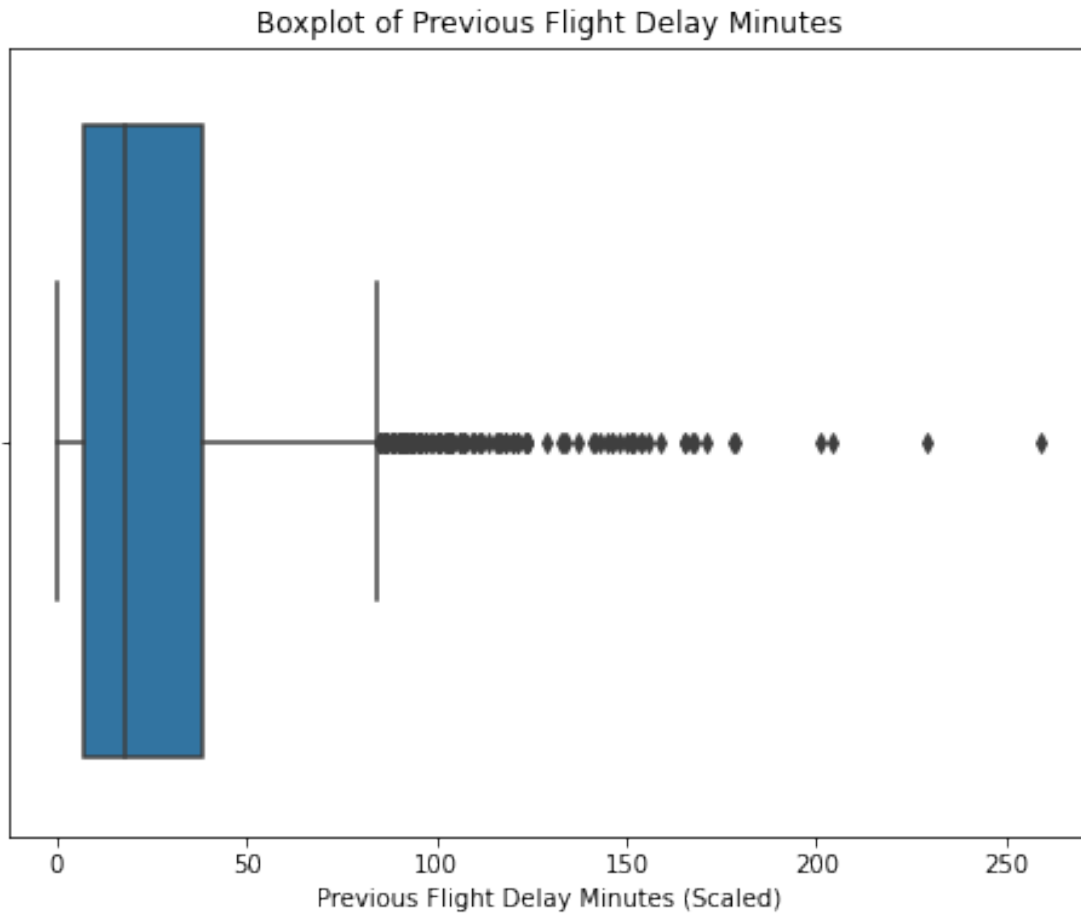
```
[23]: plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Month'])
plt.title('Boxplot of Month')
plt.xlabel('Month (Scaled)')
plt.show()
```



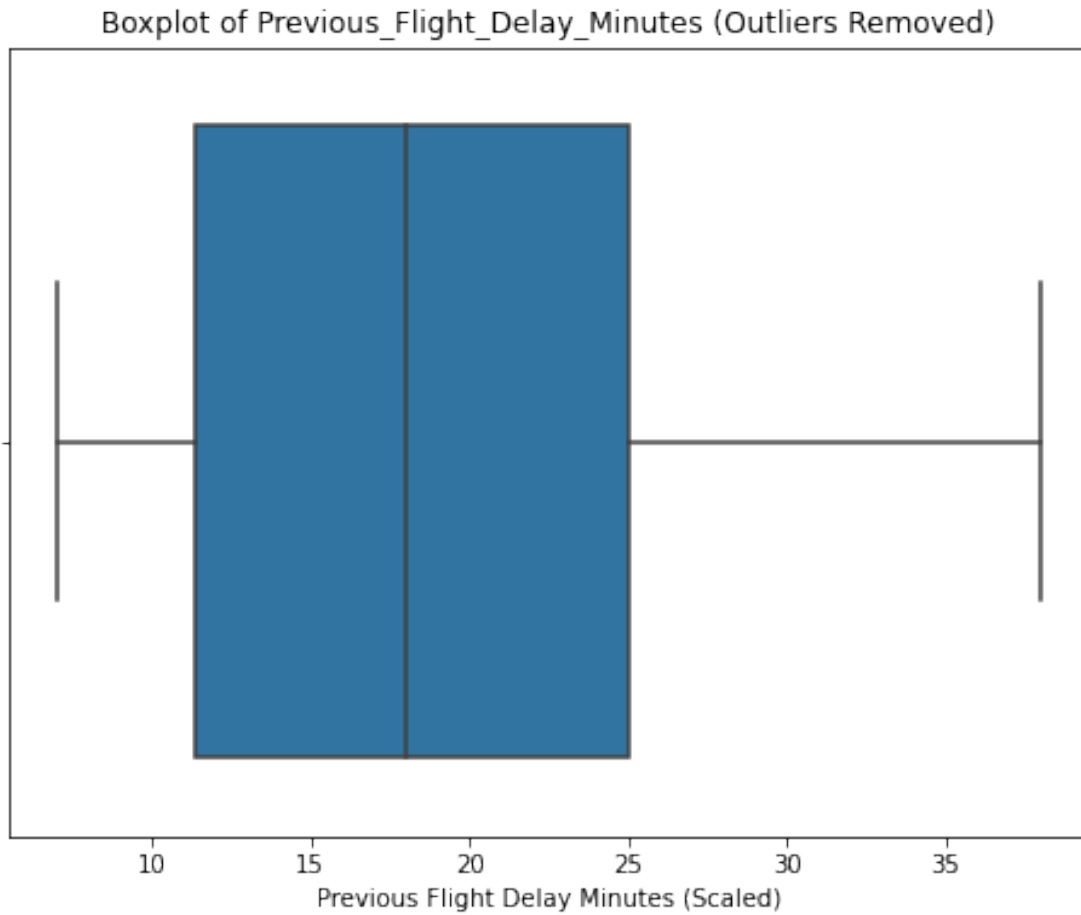
```
[24]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Weather_Score'])
plt.title('Boxplot of Weather Score')
plt.xlabel('Weather Score (Scaled)')
plt.show()
```



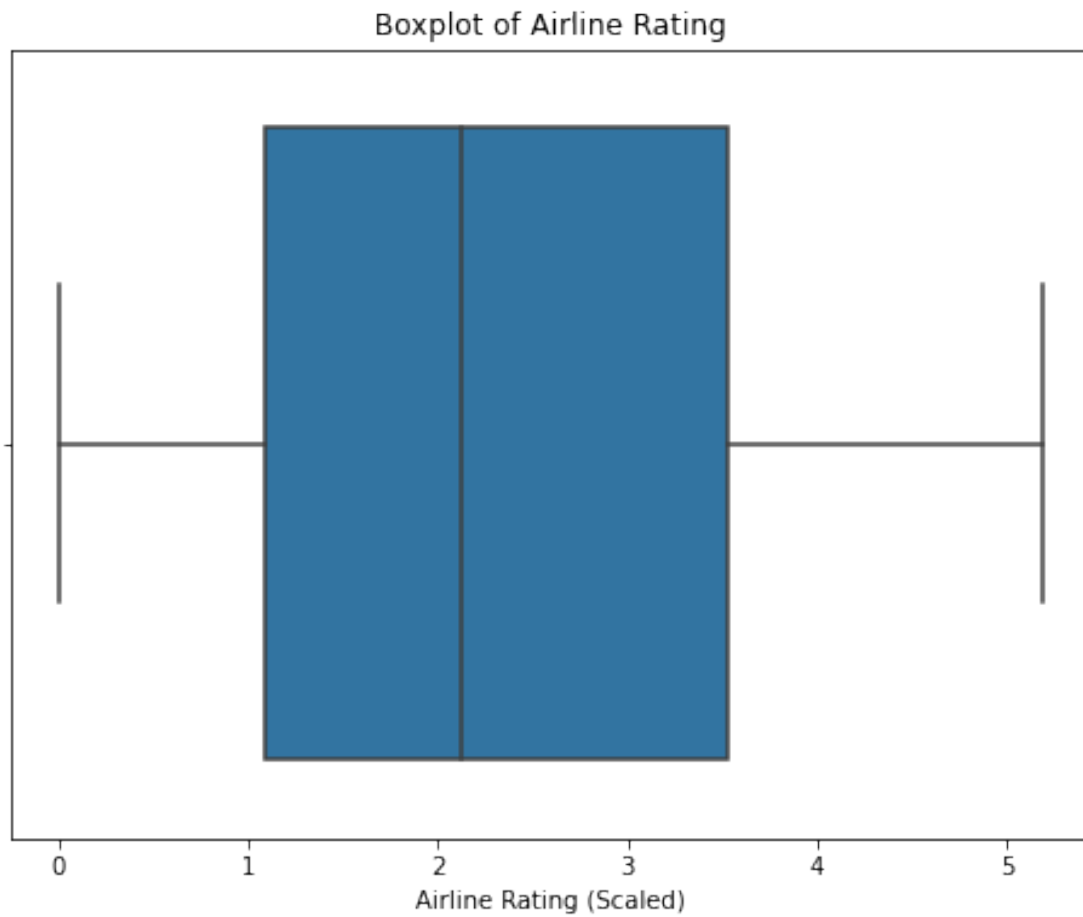
```
[25]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Previous_Flight_Delay_Minutes'])
plt.title('Boxplot of Previous Flight Delay Minutes')
plt.xlabel('Previous Flight Delay Minutes (Scaled)')
plt.show()
```



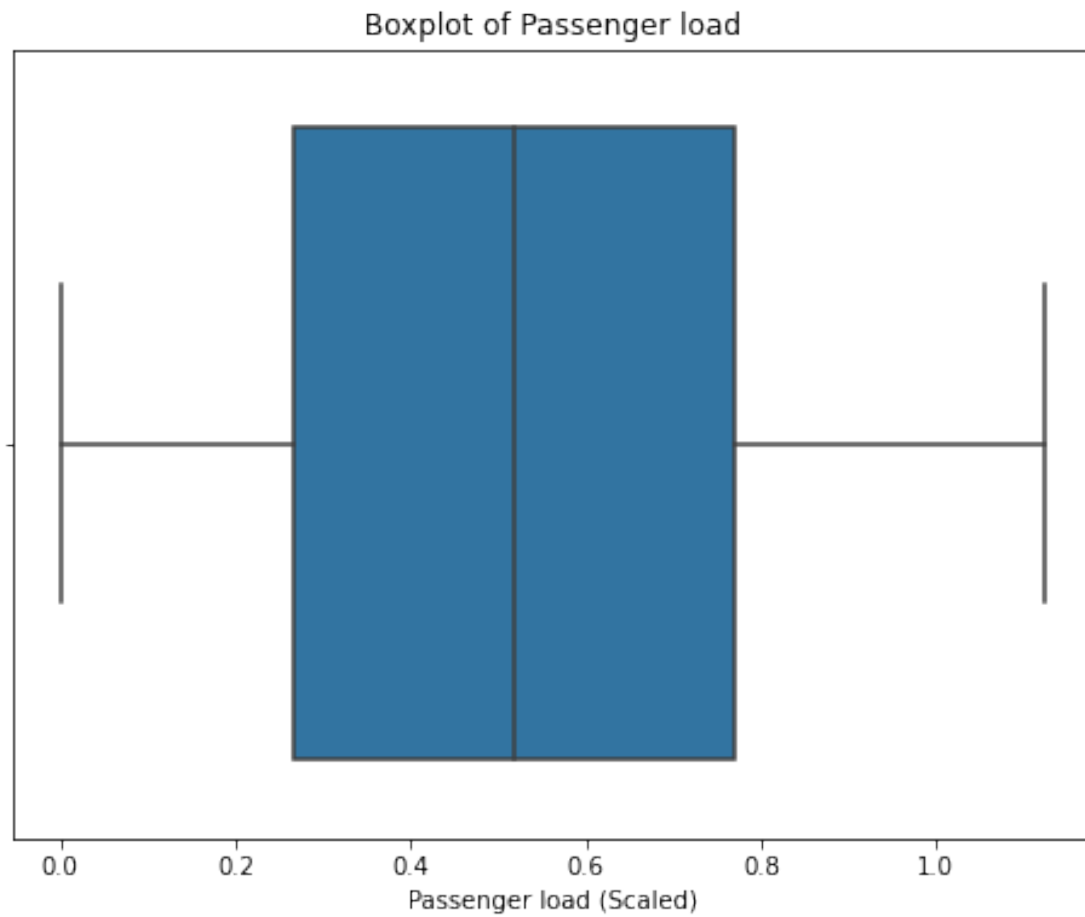
```
[26]: Q1 = Data['Previous_Flight_Delay_Minutes'].quantile(0.25)
Q3 = Data['Previous_Flight_Delay_Minutes'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 0.0 * IQR
upper_bound = Q3 + 0.0 * IQR
filtered_Data = Data[(Data['Previous_Flight_Delay_Minutes'] >= lower_bound) &
    (Data['Previous_Flight_Delay_Minutes'] <= upper_bound)]
plt.figure(figsize=(8, 6))
sns.boxplot(x=filtered_Data['Previous_Flight_Delay_Minutes'])
plt.title('Boxplot of Previous_Flight_Delay_Minutes (Outliers Removed)')
plt.xlabel('Previous Flight Delay Minutes (Scaled)')
plt.show()
```




```
[27]: plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Airline_Rating'])
plt.title('Boxplot of Airline Rating')
plt.xlabel('Airline Rating (Scaled)')
plt.show()
```

```
[28]: plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Passenger_Load'])
plt.title('Boxplot of Passenger load')
plt.xlabel('Passenger load (Scaled)')
plt.show()
```



```
[29]: plt.figure(figsize=(8, 6))
sns.boxplot(x=Data['Flight_Cancelled'])
plt.title('Boxplot of Flight Cancelled')
plt.xlabel('Flight Cancelled (Scaled)')
plt.show()
```



```
[30]: data = pd.read_csv('Flyzy Flight Cancellation - Sheet1.csv')
```

```
[31]: data.describe()
```

```
[31]:
```

	Flight ID	Flight_Distance	Scheduled_Departure_Time	Day_of_Week \
count	3.000000e+03	3000.000000	3000.000000	3000.000000
mean	4.997429e+06	498.909333	11.435000	3.963000
std	2.868139e+06	98.892266	6.899298	2.016346
min	3.681000e+03	138.000000	0.000000	1.000000
25%	2.520313e+06	431.000000	6.000000	2.000000
50%	5.073096e+06	497.000000	12.000000	4.000000
75%	7.462026e+06	566.000000	17.000000	6.000000
max	9.999011e+06	864.000000	23.000000	7.000000

	Month	Weather_Score	Previous_Flight_Delay_Minutes \
count	3000.000000	3000.000000	3000.000000
mean	6.381000	0.524023	26.793383
std	3.473979	0.290694	27.874733

min	1.000000	0.000965	0.000000
25%	3.000000	0.278011	7.000000
50%	6.000000	0.522180	18.000000
75%	9.000000	0.776323	38.000000
max	12.000000	1.099246	259.000000

	Airline_Rating	Passenger_Load	Flight_Cancelled
count	3000.000000	3000.000000	3000.000000
mean	2.317439	0.515885	0.690667
std	1.430386	0.295634	0.462296
min	0.000103	0.001039	0.000000
25%	1.092902	0.265793	0.000000
50%	2.126614	0.517175	1.000000
75%	3.525746	0.770370	1.000000
max	5.189038	1.123559	1.000000

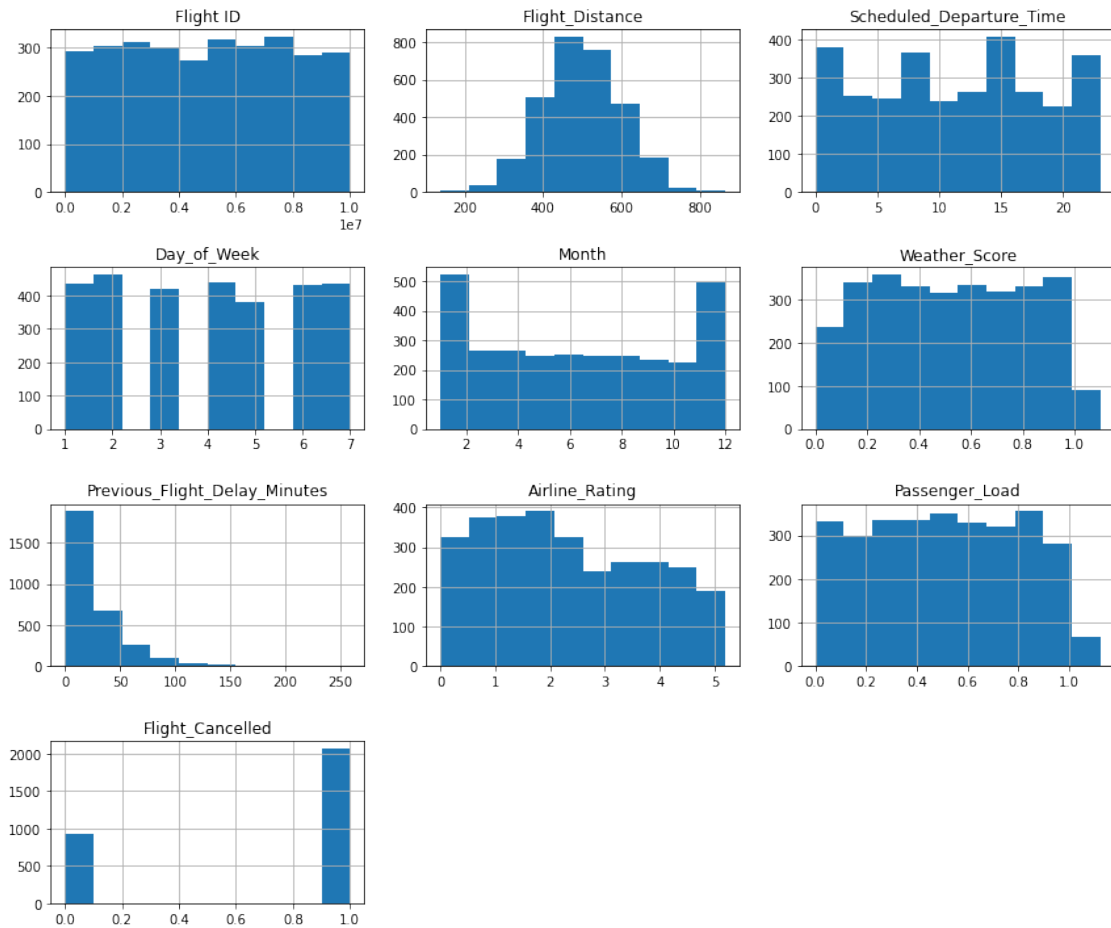
[32]: Data.info()

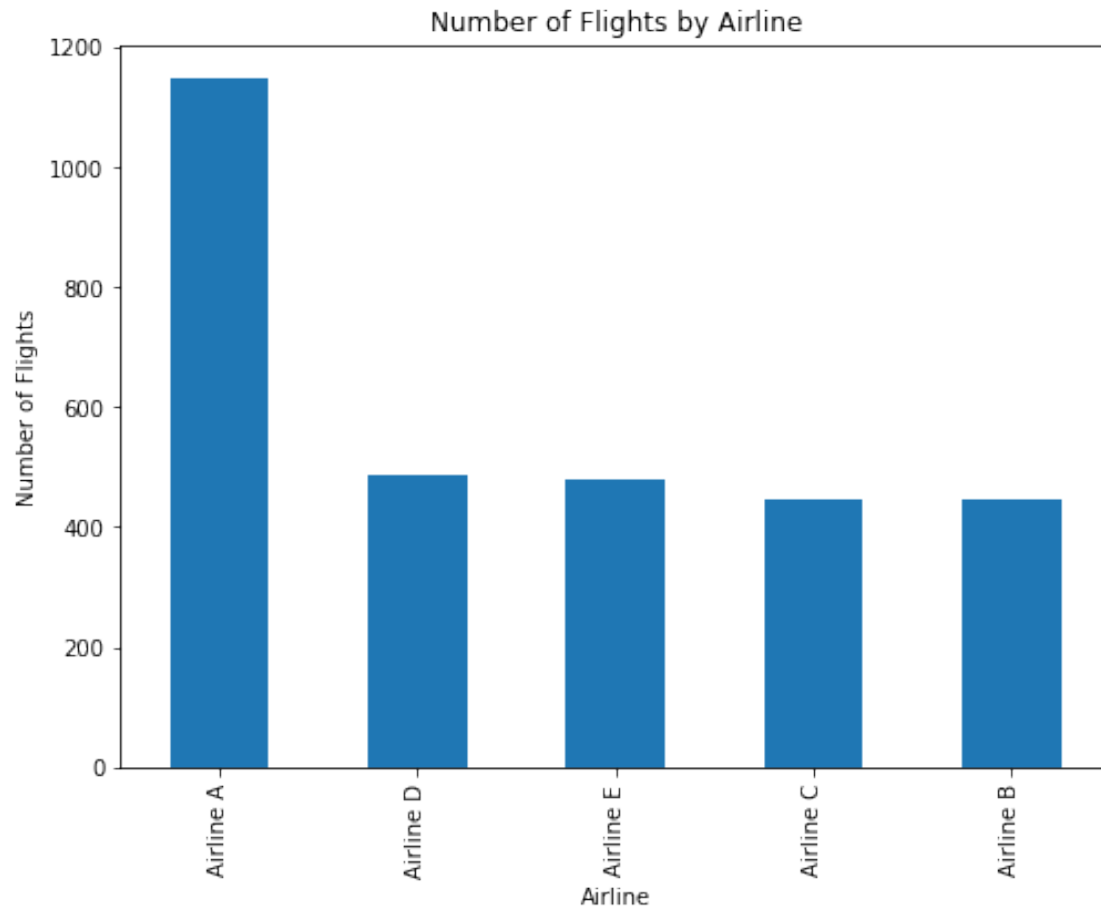
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Flight ID                            3000 non-null   int64
1   Airline                              3000 non-null   object
2   Flight_Distance                      3000 non-null   int64
3   Origin_Airport                      3000 non-null   object
4   Destination_Airport                 3000 non-null   object
5   Scheduled_Departure_Time             3000 non-null   int64
6   Day_of_Week                         3000 non-null   int64
7   Month                               3000 non-null   int64
8   Airplane_Type                       3000 non-null   object
9   Weather_Score                      3000 non-null   float64
10  Previous_Flight_Delay_Minutes        3000 non-null   float64
11  Airline_Rating                      3000 non-null   float64
12  Passenger_Load                      3000 non-null   float64
13  Flight_Cancelled                    3000 non-null   int64
dtypes: float64(4), int64(6), object(4)
memory usage: 328.2+ KB
```

[33]:


```
Data.hist(figsize=(12, 10))
plt.tight_layout()
plt.show()
plt.figure(figsize=(8, 6))
Data['Airline'].value_counts().plot(kind='bar')
plt.title('Number of Flights by Airline')
plt.xlabel('Airline')
```

```
plt.ylabel('Number of Flights')
plt.show()
print("\n")
```







```
[34]: plt.figure(figsize=(8, 6))
Data['Airline'].value_counts().plot(kind='bar')
plt.title('Number of Flights by Airline')
plt.xlabel('Airline')
plt.ylabel('Number of Flights')
plt.show()
```




```
[35]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Flight_Distance'], kde=True, bins=20, color='blue')
plt.title('Histogram of Flight Distance')
plt.xlabel('Flight Distance')
plt.ylabel('Frequency')
plt.show()
```



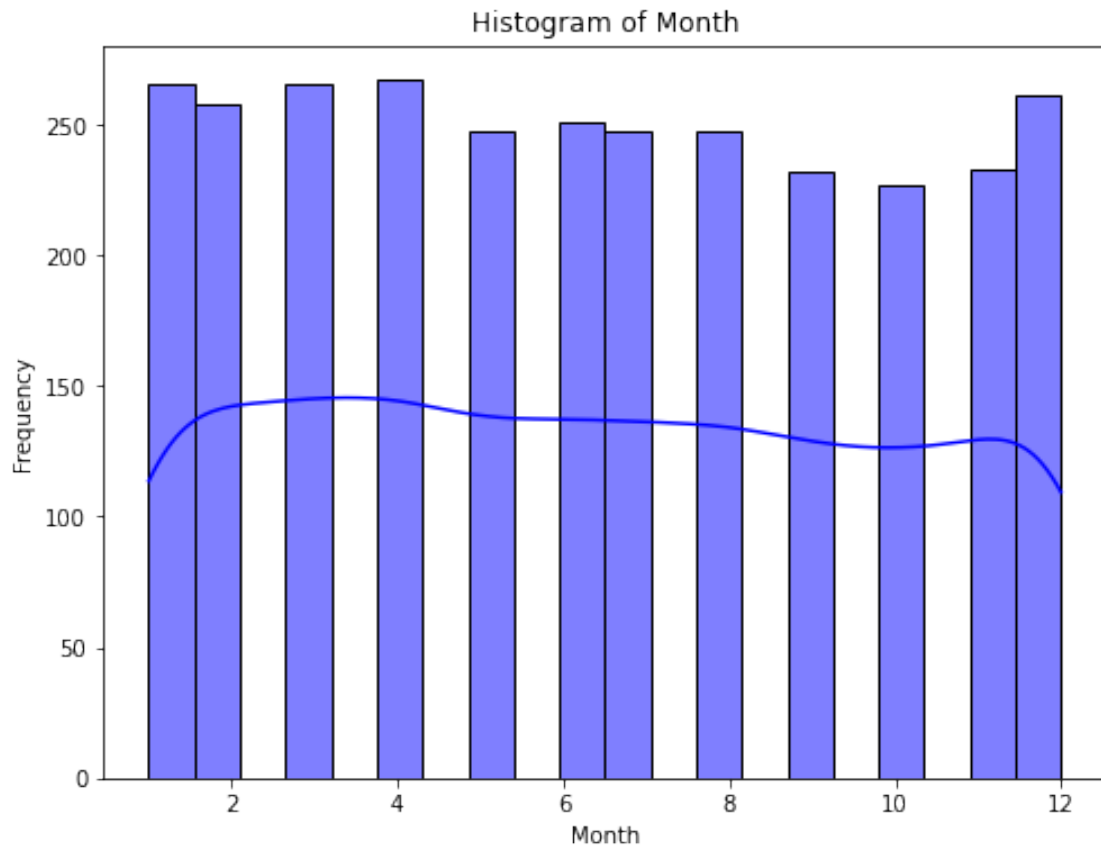
```
[36]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Scheduled_Departure_Time'], kde=True, bins=20, color='blue')
plt.title('Histogram of Scheduled Departure Time')
plt.xlabel('Scheduled Departure Time')
plt.ylabel('Frequency')
plt.show()
```


```
[37]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Day_of_Week'], kde=True, bins=20, color='blue')
plt.title('Histogram of Day of Week')
plt.xlabel('Day of Week')
plt.ylabel('Frequency')
plt.show()
```



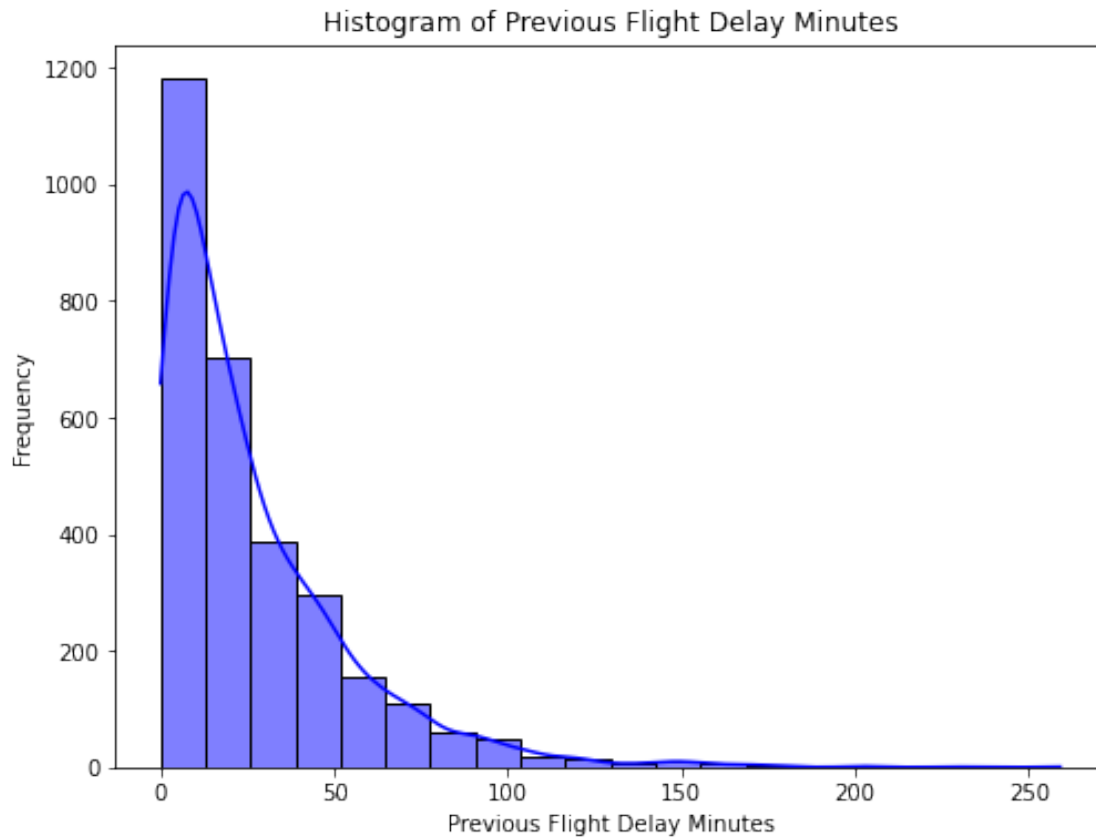
```
[38]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Month'], kde=True, bins=20, color='blue')
plt.title('Histogram of Month')
plt.xlabel('Month')
plt.ylabel('Frequency')
plt.show()
```




```
[39]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Weather_Score'], kde=True, bins=20, color='blue')
plt.title('Histogram of Weather_Score')
plt.xlabel('Weather_Score')
plt.ylabel('Frequency')
plt.show()
```




```
[40]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Previous_Flight_Delay_Minutes'], kde=True, bins=20, color='blue')
plt.title('Histogram of Previous Flight Delay Minutes')
plt.xlabel('Previous Flight Delay Minutes')
plt.ylabel('Frequency')
plt.show()
```



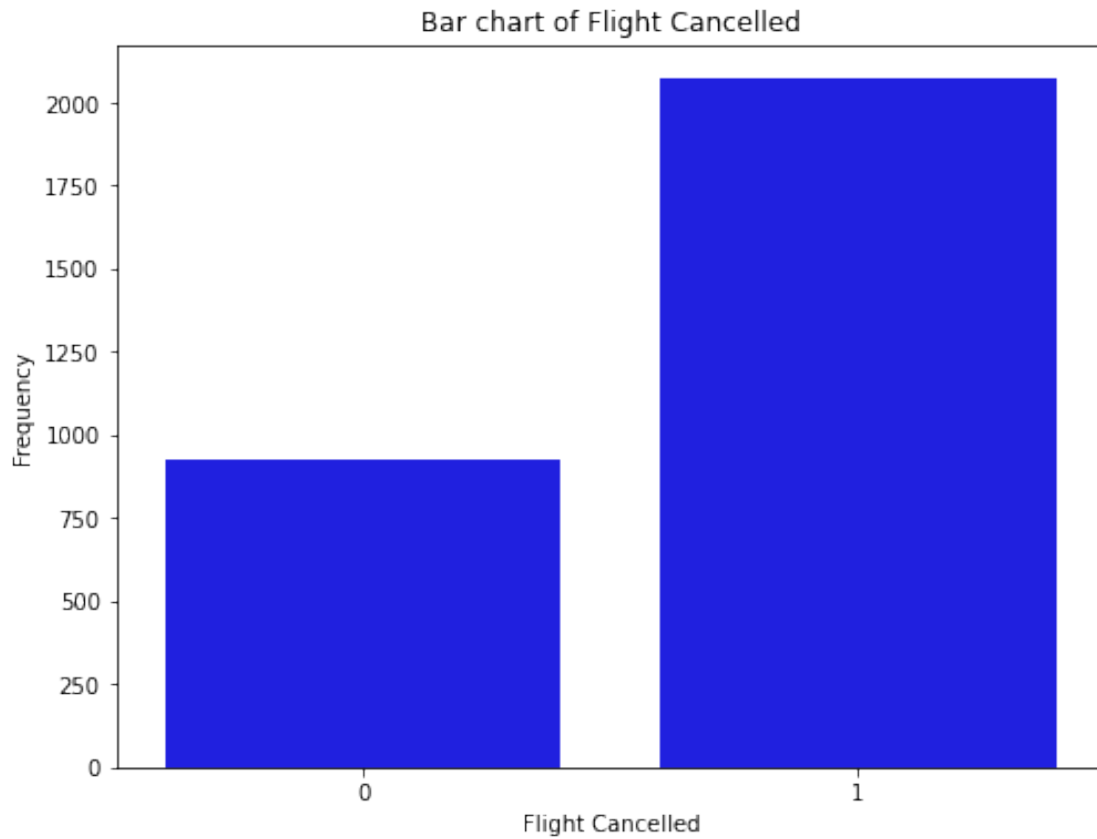
```
[41]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Airline_Rating'], kde=True, bins=20, color='blue')
plt.title('Histogram of Airline Rating')
plt.xlabel('Airline Rating')
plt.ylabel('Frequency')
plt.show()
```



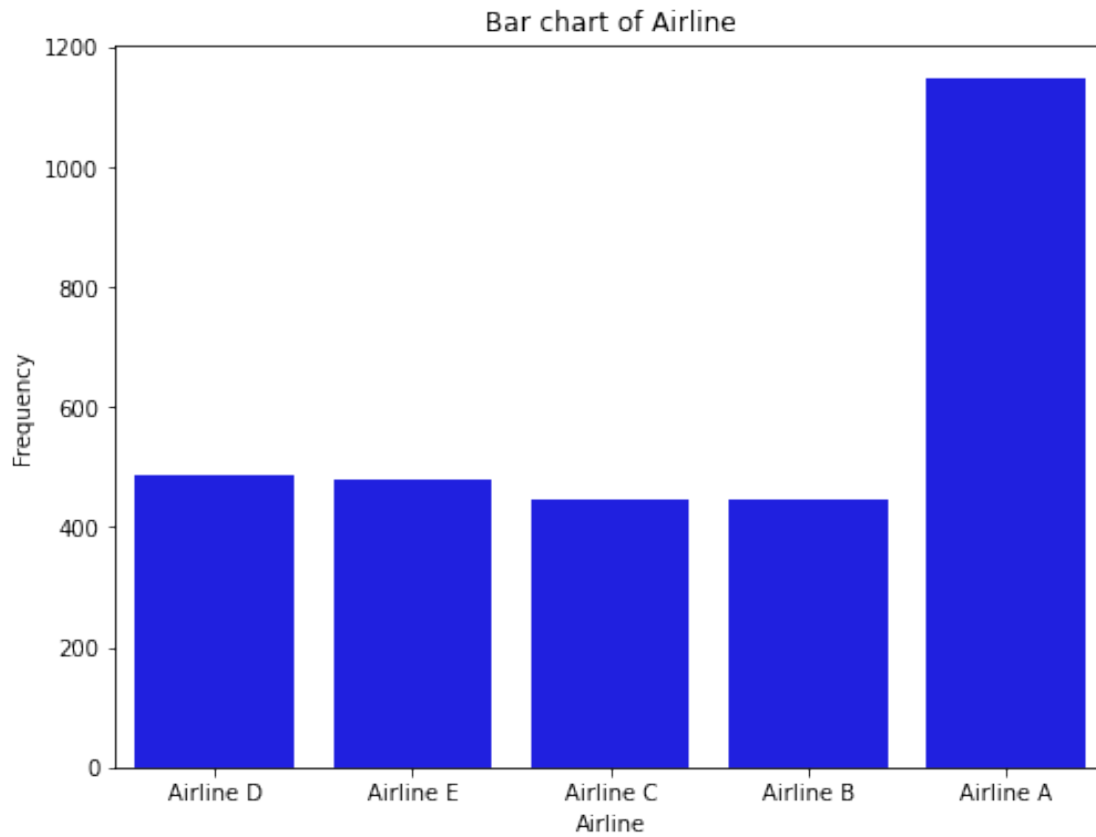
```
[42]: plt.figure(figsize=(8, 6))
sns.histplot(Data['Passenger_Load'], kde=True, bins=20, color='blue')
plt.title('Histogram of Passenger Load')
plt.xlabel('Passenger Load')
plt.ylabel('Frequency')
plt.show()
```




```
[43]: plt.figure(figsize=(8, 6))
sns.countplot(data=Data, x= 'Flight_Cancelled', color='blue')
plt.title('Bar chart of Flight Cancelled')
plt.xlabel('Flight Cancelled')
plt.ylabel('Frequency')
plt.show()
```



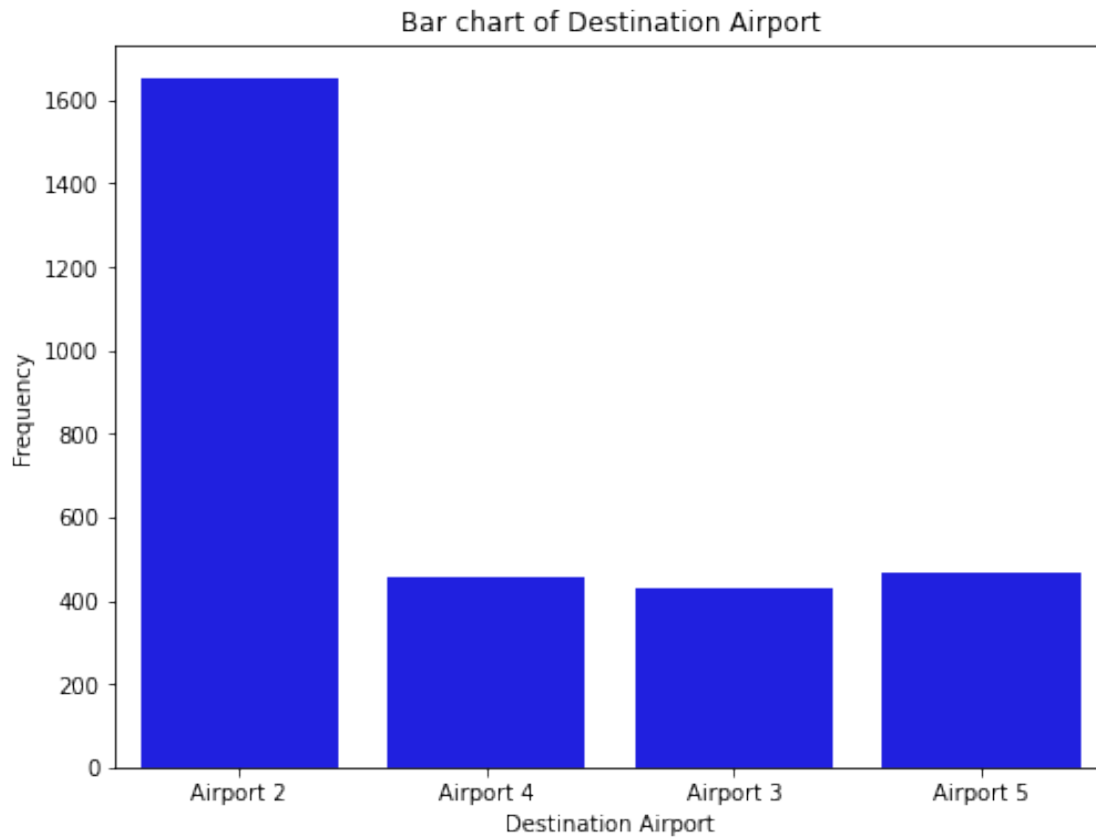
```
[44]: plt.figure(figsize=(8, 6))
sns.countplot(data=Data, x= 'Airline', color='blue')
plt.title('Bar chart of Airline')
plt.xlabel('Airline')
plt.ylabel('Frequency')
plt.show()
```

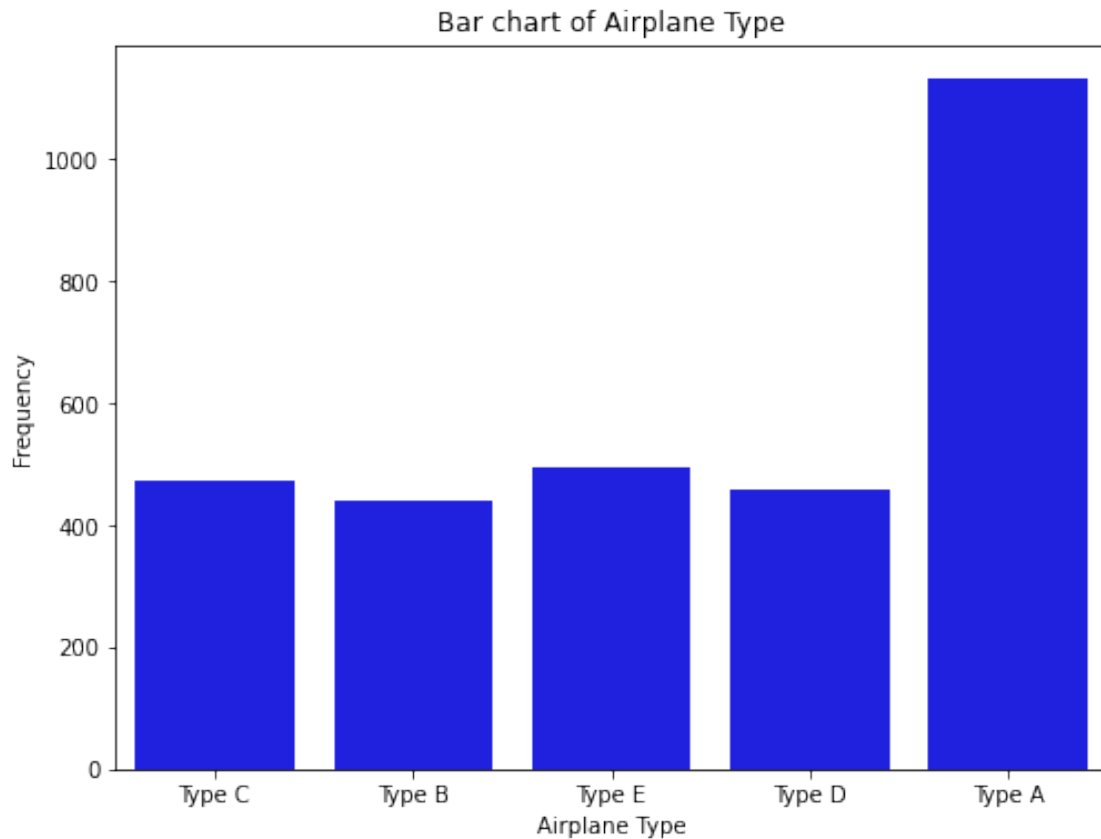
```
[45]: plt.figure(figsize=(8, 6))
sns.countplot(data=Data, x= 'Origin_Airport', color='blue')
plt.title('Bar chart of Origin Airport')
plt.xlabel('Origin Airport')
plt.ylabel('Frequency')
plt.show()
```



```
[46]: plt.figure(figsize=(8, 6))
sns.countplot(data=Data, x= 'Destination_Airport', color='blue')
plt.title('Bar chart of Destination Airport')
plt.xlabel('Destination Airport')
plt.ylabel('Frequency')
plt.show()
```



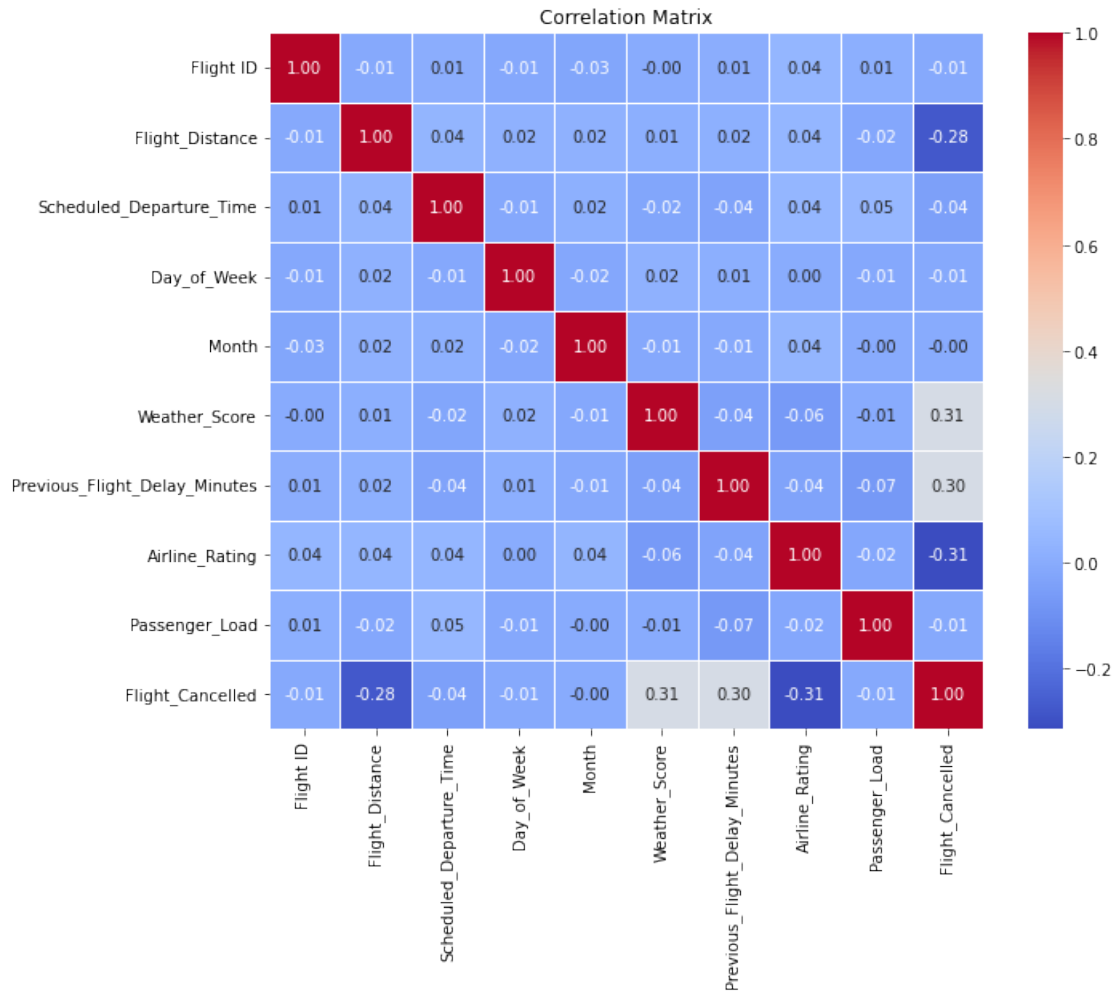
```
[47]: plt.figure(figsize=(8, 6))
sns.countplot(data=Data, x= 'Airplane_Type', color='blue')
plt.title('Bar chart of Airplane Type')
plt.xlabel('Airplane Type')
plt.ylabel('Frequency')
plt.show()
```



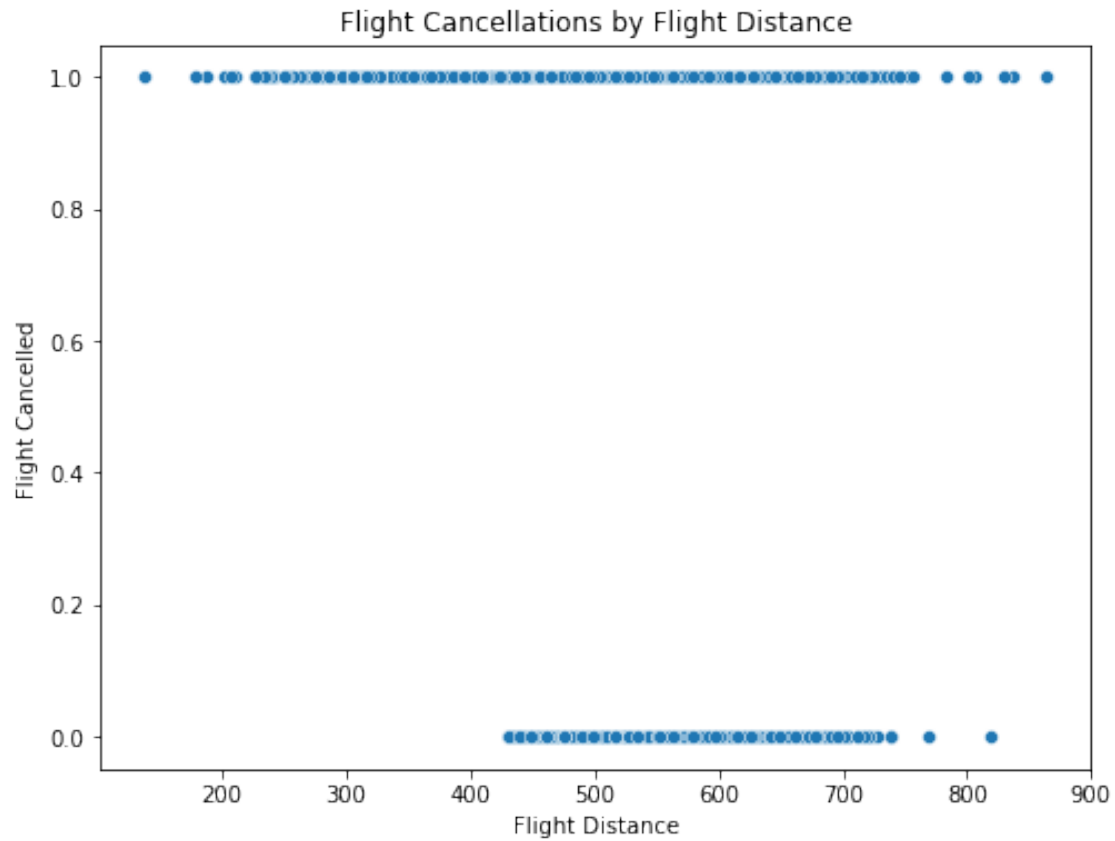
```
[48]: plt.figure(figsize=(10, 8))
sns.heatmap(Data.corr(), annot=True, cmap= 'coolwarm', fmt=".2f", linewidths=0.
↪5)
plt.title('Correlation Matrix')
plt.show()
print("\n")
```

/tmp/ipykernel_194/3230157905.py:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.


```
sns.heatmap(Data.corr(), annot=True, cmap= 'coolwarm', fmt=".2f",
linewidths=0.5)
```




```
[49]: plt.figure(figsize=(8, 6))
sns.scatterplot(data=Data, x='Flight_Distance', y='Flight_Cancelled')
plt.title('Flight Cancellations by Flight Distance')
plt.xlabel('Flight Distance')
plt.ylabel('Flight Cancelled')
plt.show()
```




```
[50]: plt.figure(figsize=(12, 6))
sns.countplot(data=Data, x='Airline', hue='Flight_Cancelled')
plt.title('Flight Cancellations by Airline')
plt.xlabel('Airline')
plt.ylabel('Number of Flights')
plt.xticks(rotation=45)
plt.legend(title='Flight Cancelled', loc='upper right')
plt.show()
```




```
[51]: plt.figure(figsize=(12, 6))
sns.countplot(data=Data, x='Origin_Airport', hue='Flight_Cancelled')
plt.title('Flight Cancellations by Origin Airport')
plt.xlabel('Origin Airport')
plt.ylabel('Number of Flights')
plt.xticks(rotation=45)
plt.legend(title='Flight Cancelled', loc='upper right')
plt.show()
```




```
[52]: plt.figure(figsize=(12, 6))
sns.countplot(data=Data, x='Destination_Airport', hue='Flight_Cancelled')
plt.title('Flight Cancellations by Destination Airport')
plt.xlabel('Destination Airport')
plt.ylabel('Number of Flights')
plt.xticks(rotation=45)
plt.legend(title='Flight Cancelled', loc='upper right')
plt.show()
```



```
[53]: plt.figure(figsize=(12, 6))
sns.countplot(data=Data, x='Airplane_Type', hue='Flight_Cancelled')
plt.title('Flight Cancellations by Airplane Type')
plt.xlabel('Airplane Type')
plt.ylabel('Number of Flights')
plt.xticks(rotation=45)
plt.legend(title='Flight Cancelled', loc='upper right')
plt.show()
```

```
[54]: numeric_data = Data.select_dtypes(include='number')
correlation_with_target = numeric_data.corr()['Flight_Cancelled'].
    ↪sort_values(ascending=False)
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_with_target.to_frame(), annot=True, cmap='coolwarm',
    ↪fmt=".2f")
plt.title('Correlation of Numeric Features with Flight Cancelled')
plt.xlabel('Numeric Features')
plt.ylabel('Flight Cancelled')
plt.show()
```



0.1 TASK 3 Preprocessing and Model Building

```
[55]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, MinMaxScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, \
f1_score, roc_auc_score
```

```
[56]: X = Data.drop(['Flight ID', 'Flight_Cancelled'], axis=1) # Features
y = Data['Flight_Cancelled']
```

```
[57]: x_encoded = pd.get_dummies(X)
```

```
[58]: print(X_encoded.head())
```

	Flight_Distance	Scheduled_Departure_Time	Day_of_Week	Month	\
0	475	4	6	1	
1	538	12	1	6	
2	565	17	3	9	
3	658	1	1	8	
4	566	19	7	12	

	Weather_Score	Previous_Flight_Delay_Minutes	Airline_Rating	\
0	0.225122	5.0	2.151974	
1	0.060346	68.0	1.600779	
2	0.093920	18.0	4.406848	
3	0.656750	13.0	0.998757	
4	0.505211	4.0	3.806206	

	Passenger_Load	Airline_Airline A	Airline_Airline B	...	\
0	0.477202	0	0	...	
1	0.159718	0	0	...	
2	0.256803	0	0	...	
3	0.504077	0	0	...	
4	0.019638	0	0	...	

	Origin_Airport_Airport 5	Destination_Airport_Airport 2	\
0	0	1	
1	1	0	
2	0	1	
3	1	0	
4	0	1	

	Destination_Airport_Airport 3	Destination_Airport_Airport 4	\
0	0	0	
1	0	1	
2	0	0	
3	1	0	
4	0	0	

	Destination_Airport_Airport 5	Airplane_Type_Type A	Airplane_Type_Type B	\
0	0	0	0	
1	0	0	1	
2	0	0	0	
3	0	0	1	
4	0	0	0	

	Airplane_Type_Type C	Airplane_Type_Type D	Airplane_Type_Type E
0	1	0	0
1	0	0	0
2	1	0	0
3	0	0	0
4	0	0	1

[5 rows x 27 columns]

```
[60]: scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X_encoded)
```

```
[61]: print(X_scaled)
```

```
[[0.46418733 0.17391304 0.83333333 ... 1.          0.          0.          ]
 [0.55096419 0.52173913 0.          ... 0.          0.          0.          ]
 [0.58815427 0.73913043 0.33333333 ... 1.          0.          0.          ]
 ...
 [0.44490358 0.34782609 0.33333333 ... 0.          0.          0.          ]
 [0.44903581 0.2173913  0.66666667 ... 0.          0.          1.          ]
 [0.31818182 0.04347826 0.          ... 0.          0.          0.          ]]
```

```
[62]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
↳ random_state=42)
print("Training set - Features:", X_train.shape, "Target:", y_train.shape)
print("Test set - Features:", X_test.shape, "Target:", y_test.shape)
```

```
Training set - Features: (2400, 27) Target: (2400,)
```

```
Test set - Features: (600, 27) Target: (600,)
```

```
[61]: print(X_scaled)
```

```
[[0.46418733 0.17391304 0.83333333 ... 1.          0.          0.          ]
 [0.55096419 0.52173913 0.          ... 0.          0.          0.          ]
 [0.58815427 0.73913043 0.33333333 ... 1.          0.          0.          ]
 ...
 [0.44490358 0.34782609 0.33333333 ... 0.          0.          0.          ]
 [0.44903581 0.2173913  0.66666667 ... 0.          0.          1.          ]
 [0.31818182 0.04347826 0.          ... 0.          0.          0.          ]]
```

```
[63]: from sklearn.linear_model import LogisticRegression
```

```
[64]: logreg_model = LogisticRegression(random_state=42)
```

```
[65]: logreg_model.fit(X_train, y_train)
```

```
[65]: LogisticRegression(random_state=42)
```

```
[66]: print("Model Coefficients:", logreg_model.coef_)
```

```
Model Coefficients: [[-5.44813124  0.01029947 -0.08940905  0.16600798
 3.36691911  8.39711858
 -2.91703632  0.05163797 -0.02845402 -0.14100312  0.0460956   0.03333066
 0.09012382 -0.02711148 -0.05280328  0.01537221  0.10468177 -0.04004627
 0.05851232 -0.03401603  0.07954138 -0.10394473  0.04275748 -0.21807963
 0.02008184  0.22047041 -0.06513717]]
```

```
[67]: print("Model Intercept:", logreg_model.intercept_)
```

Model Intercept: [2.66404546]

```
[68]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
```

```
[69]: y_pred = logreg_model.predict(X_test)
```

```
[70]: accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)
```

```
[71]: print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)
print("ROC AUC Score:", roc_auc)
```

Accuracy: 0.7883333333333333
Precision: 0.8149779735682819
Recall: 0.8958837772397095
F1-score: 0.8535178777393311
ROC AUC Score: 0.7233429581385713

```
[6]: data = pd.read_csv('Flyzy Flight Cancellation - Sheet1.csv')
```

```
[8]: X = Data.drop(columns=['Scheduled_Departure_Time', 'Flight_Cancelled'])
y = Data['Flight_Cancelled']
numerical_features = ['Flight_Distance', 'Scheduled_Departure_Time',
    'Previous_Flight_Delay_Minutes', 'Airline_Rating', 'Passenger_Load',
    'Weather_Score']
categorical_features = ['Airline', 'Origin_Airport', 'Destination_Airport',
    'Airplane_Type', 'Day_of_Week', 'Month']
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numerical_features),
        ('cat', OneHotEncoder(), categorical_features)
    ])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)
```

```
-----
NameError                                Traceback (most recent call last)
/tmp/ipykernel_72/185815045.py in <cell line: 1>()
----> 1 X = Data.drop(columns=['Scheduled_Departure_Time', 'Flight_Cancelled'])
      2 y = Data['Flight_Cancelled']
```

```
3 numerical_features = ['Flight_Distance', 'Scheduled_Departure_Time',  
↪ 'Previous_Flight_Delay_Minutes', 'Airline_Rating', 'Passenger_Load',  
↪ 'Weather_Score']  
4 categorical_features = ['Airline', 'Origin_Airport',  
↪ 'Destination_Airport', 'Airplane_Type', 'Day_of_Week', 'Month']  
5 preprocessor = ColumnTransformer()
```

NameError: name 'Data' is not defined

[]: